

A C Data Structure Book

Matej Dedina

October 8, 2026

Preface

This book describes the process of designing and implementing a personal data structure library in the C programming language. Each implementation is centered around two types of implementations - an infinitely expandable structure able to hold any amount of elements, and a finite one able to hold only a certain maximum amount.

The goal of this work is not to be the go-to data structure book for C, on the contrary, the name is meant to represent an affordable book for the public which doesn't take itself seriously.

The book will be divided into parts classifying all structures into categories of chapters. Each chapter will explain the process of creating said structure and the code used to make it.

Since the code is in C a decent understanding of the programming language is required, that primarily includes - structures (struct), pointers, memory allocation, and last but certainly not least - function pointers.

All the source code will be available at github under the Unlicense License.

If I have to put a disclaimer I just want to clarify that the reason why some paragraphs may be written like this is because the book was created using $\text{\LaTeX}$ in Visual Studio Code and I hate seeing `Overfull \hbox` highlights.

Contents

1	The cerpec data structure library	1
1.1	The main header	1
1.1.1	Growth factor and chunk size	1
1.1.2	Error and validation handling	2
1.1.3	Custom memory allocator	3
1.1.4	Function pointers	5
1.2	Implementation philosophies	6
1.2.1	Mantras	6
1.2.2	How others did it	7
1.2.3	Python	7
1.2.4	Haskell	8
1.3	On a final note...	8
	Further reading	8
I	Restricted sequential data structures	9
2	Stack	13
2.1	Implementations of a Stack	16
2.1.1	The humble linked list	16
2.1.2	The prideful array	16
2.1.3	Singly linked list of array elements	17
2.2	Array stack structure code	18
2.2.1	The infinite stack structure	18
2.2.2	Resize stack	19
2.2.3	Create stack	20
2.2.4	Make stack	20
2.2.5	Destroy stack	21
2.2.6	Clear stack	22
2.2.7	Copy stack	23
2.2.8	Is empty stack	25

2.2.9	Push to stack	25
2.2.10	Pop from stack	26
2.2.11	Peep the stack	27
2.2.12	For each element	27
2.2.13	Apply to all elements	28
2.3	Example usages	29
2.3.1	Creating clearing and destroying a stack	29
2.3.2	Pushing, peeping and popping elements	29
2.3.3	Iterating and sorting elements	30
2.4	On a final note...	32
	Further reading	32
3	Queue	33
3.1	Implementations of a queue	33
3.1.1	The free linked list	35
3.1.2	The confined array	35
3.1.3	Tail-based circular linked list of arrays	37
3.2	Tail-based circular linked list of array code	37
3.2.1	The infinite queue structure	37
3.2.2	Create queue	39
3.2.3	Make queue	40
3.2.4	Destroy queue	40
3.2.5	Clear queue	42
3.2.6	Copy queue	43
3.2.7	Is queue empty	45
3.2.8	Enqueue	46
3.2.9	Dequeue	47
3.2.10	Peek queue	48
3.2.11	For each element	49
3.2.12	Apply to all elements	50
3.3	Example usages	52
3.3.1	Creating, clearing and destroying a queue	52
3.3.2	Enqueueing, peeking and dequeueing elements	53
3.3.3	Iterating and sorting elements	54
3.4	On a final note...	55
	Further reading	55
4	Deque	57
4.1	Implementations of a deque	57
4.1.1	The simple array	58
4.1.2	The complex doubly-linked list	59

4.1.3	Doubly-linked list of arrays	60
4.2	Doubly-linked list of arrays code	60
4.2.1	The infinite queue structure	60
4.2.2	Create deque	61
4.2.3	Make deque	61
4.2.4	Destroy deque	62
4.2.5	Clear deque	63
4.2.6	Copy deque	64
4.2.7	Is empty deque	67
4.2.8	Enqueue deque front	67
4.2.9	Enqueue deque back	69
4.2.10	Dequeue deque front	71
4.2.11	Dequeue deque back	72
4.2.12	Peek deque front	74
4.2.13	Peek deque back	74
4.2.14	Each deque front	75
4.2.15	Each deque back	76
4.2.16	Apply deque	77
4.3	Example usages	80
4.3.1	Creating, cleaning and destroying a deque	80
4.3.2	Enqueueing, peeking and dequeuing elements front and back	80
4.3.3	Iterating and sorting elements	81
4.4	On a final note...	83
	Further reading	83

II Linked lists 85

5 Straight list 89

5.1	Implementations of a straight list	89
5.1.1	A bread of crumbs - the linked list	91
5.1.2	A slice-off emmental - the sparse array	91
5.1.3	A stack of butter - the straight list	94

Chapter 1

The cerpec data structure library

The name `cerpec` may look familiar if you're a Polish speaker, but the word actually comes from the Eastern Slovak word for suffering. One can also see the similarities when writing and reading the standard Slovak, Eastern Slovak, and Polish words next to each other, like this `trpieť - cerpec - cierpieć`.

The reasons for choosing this name are three-fold; making this library was like suffering through psychological pain; Eastern Slovak infinitives ending with `-c` (standard Slovak with `-ť`) go nicely with the `C` in the C programming language; and lastly, I doubt anyone would use `cerpec` for trademarking reasons since it's a dialectic word.

1.1 The main header

The `cerpec.h` serves as the parent containing all the necessary definitions shared among all the data structures. The header can be divided into four parts.

1.1.1 Growth factor and chunk size

The growth factor `CERPEC_FACTOR` represents the factor by which to determine the previous and next memory size to resize into when we're either out of space or if smaller space is available.

```
include/cerpec.h
```

```
#define CERPEC_FACTOR 2
```

The chunk size `CERPEC_CHUNK` is used to avoid a small starting size where the program might be too busy just reallocating the structures instead of performing operations, graphically exemplified in Figure 1.1. Basically, instead of starting at length 1 after inserting the first element, it starts at a specific higher power of two, in the case of this project - 256, and grows exponentially as in Listing 1.1.1.




Figure 1.1: Example growth if chunk size is 1 (top) vs 8 (bottom). The height of the pyramid represents the number of memory growths performed.

include/cerpec.h

```
// define chunk size to expand and contract all data structures
#if !defined(CERPEC_CHUNK)
#   define CERPEC_CHUNK 256
#elif CERPEC_CHUNK ≤ 0
#   error "Chunk size must be greater than zero."
#elif (CERPEC_CHUNK & (CERPEC_CHUNK - 1))
#   error "Chunk size must be a power of 2."
#endif
```

1.1.2 Error and validation handling

Assertions are a simple way to check if certain properties are true either before or after an action has been performed. For cerpec there are two types of assertions - error and valid (see Listing 1.1.2).

The error assertion errors when something bad that shouldn't happen happens, common examples include - a parameter being NULL, memory allocation failing, no element being found. The valid assertion validates the state of the structure, for example if length member is less than or equal to capacity member, if element size member is greater than zero. valid assertions are just less important conditions that should be

disabled in production for better performance.

It's just a means to allow the user to specify which assertions they want to check and disable. The way to disable `valid`, `error` and all assertions altogether are via defining `NVALID`, `NERROR` and `NDEBUG` respectively, either through compiler flag macro definitions or using `#define` prior to any data structure inclusion. I believe that in the future more assertions will be added.

`include/cerpec.h`

```
#ifndef NDEBUG
#   define NVALID
#   define NERROR
#endif

#ifndef NVALID
#   define valid(condition) assert(condition)
#else
#   define valid(condition) (void)(0)
#endif

#ifndef NERROR
#   define error(condition) assert(condition)
#else
#   define error(condition) (void)(0)
#endif
```

1.1.3 Custom memory allocator

Each structure has a pointer to a custom memory allocator which can be used to allocate, free and reallocate memory based on arguments. The standard memory allocator is a constant external variable that uses the standard library's `malloc`, `free` and `realloc` functions in the `create_*` data structure functions. Below is an example of how the book shows standard C library functions, these ones taken from the Linux manual pages.

Linux manual page - `malloc(3)`

The `malloc()` function allocates `size` bytes and returns a pointer to the allocated memory. The memory is not initialized. If `size` is 0, then `malloc()` returns a unique pointer value that can later be successfully passed to `free()`.

Linux manual page - free(3)

The `free()` function frees the memory space pointed to by `p`, which must have been returned by a previous call to `malloc()` or related functions. Otherwise, or if `p` has already been freed, undefined behavior occurs. If `p` is `NULL`, no operation is performed.

Linux manual page - realloc(3)

The `realloc()` function changes the size of the memory block pointed to by `p` to `size` bytes. The contents of the memory will be unchanged in the range from the start of the region up to the minimum of the old and new sizes. If the new size is larger than the old size, the added memory will not be initialized. If `p` is `NULL`, then the call is equivalent to `malloc(size)`, for all values of `size`. If `size` is equal to zero, and `p` is not `NULL`, then the call is equivalent to `free(p)`. Unless `p` is `NULL`, it must have been returned by an earlier call to `malloc` or related functions. If the area pointed to was moved, a `free(p)` is done.

The only way to use custom memory is through defining three custom function pointers and putting them as parameters for the `compose_memory` builder, as is the case in Listing 1.1.3. Each data structure has a function starting with `make_` which allows more customization. Right now it only allows to add a custom allocator, but in the future it might allow more parameters for customization.

```
include/cerpec.h
```

```
typedef void * (*alloc_fn)    (size_t const, void *);
typedef void * (*realloc_fn) (void *, size_t const, void *);
typedef void   (*free_fn)     (void *, void *);

typedef struct memory {
    void * arguments;
    alloc_fn alloc;
    realloc_fn realloc;
    free_fn free;
} memory_s;

memory_s compose_memory(alloc_fn const alloc, realloc_fn const realloc,
    free_fn const free, void * const arguments);
```

The standard memory allocations can be accessed through the constant external variable `standard`.

```
include/cerpec.h
```

```
extern const memory_s standard;
```

1.1.4 Function pointers

Listing 1.1.4 is a code snippet list of all available function pointers which allow for more generic and user-defined data manipulation. One minor philosophy I would like to point out is the usage of `void*` arguments for the function pointers. This concept is based on reentrant-ness, specifically from the `qsort_r` function in `stdlib.h`. In it, the compare function pointer for `qsort` uses a third argument which allows it to avoid using global variables and thus allows the user to use threads more safely. From my personal perspective, I don't care about multithreading, but I needed those reentrant function pointers for my other personal ECS project.

Linux manual page - `qsort_r(3)`

The `qsort_r()` function is identical to `qsort()` except that the comparison function `compar` takes a third argument. A pointer is passed to the comparison function via `arg`. In this way, the comparison function does not need to use global variables to pass through arbitrary arguments, and is therefore reentrant and safe to use in threads.

```
include/cerpec.h
```

```
typedef void (*set_fn) (void * const element, void * arg);
typedef void * (*copy_fn) (void * const destination, void const *
    const source, void * arg);
typedef size_t (*hash_fn) (void const * const element, void * arg);
typedef int (*compare_fn) (void const * const a, void const * const
    b, void * arg);
typedef bool (*filter_fn) (void const * const element, void * arg);
typedef bool (*manage_fn) (void * const element, void * arg);
typedef void (*process_fn) (void * const array, size_t const lenght,
    void * arg);
typedef void (*operate_fn) (void * const result, void const * const
    a, void const * const b, void * arg);
```

The explanation for each of the function pointers is as follows:

1. `set_fn` - a function pointer that allows to manipulate a generic element pointer directly, can be used to set the specified element to zero, deallocate memory, increment and decrement values, etc.

2. `copy_fn` - a function pointer to create a copy of a specific element. It is possible to also copy all nested sub-elements either deeply or referenced shallowly, starting from a source element reference into a destination one.
3. `hash_fn` - a function pointer which takes in a generic element and returns a `size_t` value representing a hash. This hash can be used to index the value into a hash array.
4. `compare_fn` - a function pointer that compares elements `a` and `b`, and returns 0 if equal, negative if `a` is smaller, and positive otherwise.
5. `filter_fn` - a function pointer that returns `true` if the specified element meets a certain criteria, `false` otherwise. Can be used to extract elements in linear list structures, for example odd/even numbers, prime numbers, but also whitespace, alphanumeric characters, etc.
6. `manage_fn` - a function pointer that can manage a single element based on arguments. This function allows to iteratively change elements until `false` is returned (kinda like a `break` in a classic `c` loop.). The closest thing to this would be a `foreach` loop body in other languages like `C#`.
7. `process_fn` - a function pointer that allows to process an array of elements instead of iteratively walking through each one. Can be used to sort a sequence of element in array form for linear data structures like stacks, queues, dequeues and lists.
8. `operate_fn` - a function pointer that performs an operation on two elements and saves the result inside the first `result` parameter. The only usage is in the $F(n) = G(n) + H(n)$ calculation for the A^* path finding algorithm ($F(n)$ being `result` and the rest being the other two parameters).

1.2 Implementation philosophies

This section simply contains my personal preferences and opinions about certain design and code structure choices. I hereby give you permission to disagree with everything written here.

1.2.1 Mantras

I made up these following mantras used by the data structures:

- Snake case.

- Public header functions end with the structure name, private source ones begin with an underscore and the name, like this public: `push_istack`, private: `_istack_resize`.
- `char * elements` - All contained element arrays must be pointers to characters. The reason is simple, the `sizeof` operator returns the number of `char` elements a type can hold. This is important for traversing the array, as getting to an element uses this formula: $element = elements_array + (index_of_element * size_of_element)$.
- East const - The code uses east const style for better readability. Since I decided to change from West const to east const while being mostly done with the library there might be some West const left. Here's west const for constant integer and constant pointer to constant integer - `const int`, `const int * const`; and east const - `int const`, `int const * const`.
- Index-relativity - Technically every array, like `element`, `next`, `prev` in lists, can be either written inside a node struct or divided into arrays where the node itself becomes an array index. Classic nodes are just indexes in the array of available memory, meanwhile dividing node members into individual arrays makes the array index relative to its data structure. Simply said, instead of having a doubly linked list with a node struct member with `element`, `next` and `prev` members, we have just the list structure with an `elements` array, `nexts` array, and `prevs` array member and the 'node' is just an index value.

1.2.2 How others did it

A short description of how other programming languages implemented certain structures and how I took creative liberty in reimplementing them.

C++

The C++ programming language and its standard data structures have been a great source of inspiration for this project. Many structures follow the same design and implementations as the C++ standard library. Every time I wasn't sure how to implement something I looked up how C++ did it and adapted it to C. With that said, there were some gaps.

1.2.3 Python

From Python I took the idea of set operations. I decided to implement them just like Monty Python intended. Jokes aside, I find the Python set operations way more human than what C++ offers.

1.2.4 Haskell

Haskell is probably the most quintessential programming language for my library. Using function pointers as parameters is the thing that defines this project and without Haskell's higher order functions serving as inspiration I wouldn't be able to make the project into reality. Haskell lists also played a significant role of inspiration. As is evident in the many operations one can perform on cerpec's lists.

1.3 On a final note...

From a more in-depth explanation of the C programming language I recommend reading [2] as this is the goto C book. Since reading a book can be exhausting and most of the stuff in it will be easily forgotten the Linux Manual Pages [3] are a quick way to look up official documentation. For a more casual explanation of the benefits of using assert in your codebase you can watch this short YouTube video [4]. Since I mainly talked about C++ and Haskell as my main sources of inspiration for many of the data structures in the next parts chapters I will also cite their references respectively and in order, [1, 5].

Further reading

- [1] cplusplus.com. C++ Reference. Accessed: 2026-03-16. 2026. URL: <https://cplusplus.com/reference/stl/>.
- [2] B.W. Kernighan and D. Ritchie. C Programming Language. Pearson Education, 1988. ISBN: 9780133086218. URL: <https://books.google.sk/books?id=Yi5FI5QcdmYC>.
- [3] Michael Kerrisk. Linux man pages online. <https://man7.org>. Accessed: 2026-03-16. Linux Foundation, 2026.
- [4] Lewboski. Code Reliability with Assertions. www.youtube.com/watch?v=ayfw_NFBWQg&t. Accessed: 2026-03-16. Nov. 2025.
- [5] NoRedInk. Module List. <https://hackage.haskell.org>. Accessed: 2026-03-16. 2026.

Part I

Restricted sequential data structures

The name of the first part comes from the inability to find a categoric name for only stacks, queues and dequeues. Online sources categorize those as linear dynamic data structures, the problem is... linked lists are also part of this category. I want to put list implementations into a separate book part, thus I hereby coin the term Restricted sequential data structures as a subcategory of dynamic linear data structures. I know nobody will take this declaration seriously, but for the sake of this book please bear with it.

Chapter 2

Stack

So, just imagine we have a stack of books. The normal way we can add books to the stack is by putting each one to the top, if we put a book with the cover pointing upwards, one also knows what exact book is at the top. To remove a book, just grab the top one and put it somewhere else.

By making space for another stack of books and adding all the books from the first stack, while the constraints of putting and removing still remain, we get a new stack, but with a specific property that, when compared with the initial one, makes us tilt our head (literally). The new stack is just like the last one, but with the books in opposite order. This little property allows us to invert the order of any linear ordering of elements without knowing how they're implemented, by just adding and removing elements.

Continuing with the two-stack analogy, what if, instead of books we stored a game of chess. By adding these moves one-by-one (see Figure 2.1): E2-E4, E7-E5, G1-F3, B8-C6, D2-D4, E5-D4, F3-D4, F8-C5, C2-C3, D8-F6, D4-C6, F6-F2; we get probably the funniest chess game in modern history ¹. And on top of that, if we remove the two moves from the top and reset the pieces (F6-F2 becomes F2-F6), iteratively, it is possible to return to the initial state of the game, see Figure 2.2.

Let me not forget to mention this - probably every data structure book and article about stacks describing this mechanism also mentions that search engines allow you to return to the previous pages, including back and forth traversal, with stacks alone.

Stacks are also used to store function calls in programming languages like C, to be able move to the previous place in memory the program should go to after the top function finishes.

¹that can be watched via <https://www.youtube.com/watch?v=e91M0XLX7Jw>




Figure 2.1: The entire game of chess between content creators xQc and MoistCr1tikal with the latter coming on top.




Figure 2.2: Visual example of two stacks being used to track the chess game, (left-right) plays the game backwards, (right-left) plays the game forward



Figure 2.3: A simple graphical linked list example.

2.1 Implementations of a Stack

Every computer science student learns that the two most common ways to implement a stack in C are:

- Singly linked lists
- Static/dynamic arrays

So why does this book have three subsections for stack implementations?

2.1.1 The humble linked list

For a more in-depth explanation of lists you can go visit and Part II's list chapter, but in short, a linked list is a linear data structure, which allows us to store data sequentially as well as store a reference to the next data item, like in Figure 2.3.

The linked list consists of two special, nodes which represent the beginning and the end - the head and the tail. Most linked list implementations will store the reference to the head, so it is generally easy to access the first element as oppose to the tail, where the user needs to follow all the arrows until they reach it.

Adding, removing and accessing the first (top of the stack) node only requires a reference to the head, as shown in Figure 2.4. Adding is just creating a new node, then making its reference point to the stack list's head, and lastly changing the head itself into the new node. Removal is storing the top node, then changing stack's head to the next reference (head's next node), and destroying the removed node.

2.1.2 The prideful array

The linked list is kinda complicated, and compared to our initial book analogy there's a better way to represent stacks - through arrays. Another problem is cache locality, which makes linked list stacks slower since elements may be stored huge memory distances away and loading them to faster, but smaller, memory areas for speed may be too slow. On the other hand, there are some improvements that can be made even with lists, but back to the topic.

An array is just a linear data structure made up of continuous memory. This simple concept makes it one of the most versatile data structures in all of programming.




Figure 2.4: Pushing element 'H' to the top of a linked list stack and creating 'HI' (left), then removing 'H' (right).

Adding, removing and accessing the first top node only requires a numerical index value that can be inferred through the stack's length. We just store the length of the stack, then add an element by putting it into the $index = length$ and incrementing length by one; removing an element requires decrementing length by one and making $index = length$. Just don't forget to check if a stack's maximum length is less than maximum/capacity size during the push; and length not being zero while popping. Accessing the topmost element is just $index = length - 1$, without the decrement (see Figure 2.5).

The biggest concern is that arrays aren't infinite and if we want to add an arbitrary number of elements we need to expand it - if we have space then we simply increase the capacity value, one must also store a changing capacity value for dynamically growing arrays and a constant maximum for static ones; else if the array occupies the maximum available chunk of memory all elements must be moved to a bigger chunk; and when that fails, because no such chunk exists, everything's screwed and we must terminate.

2.1.3 Singly linked list of array elements

Combine the best of humbleness with pride and you get this. It's just a linked list, see Figure 2.6, but the node is actually made up of an array of elements instead of just a singleton. A mathematical person may say that the previous implementations are but a simple subset of this superset of THE computer science stack, but i digress.

The main benefit is that the list-array works like an array when accessing elements; and like a list when shortening/expanding it (we don't need to resize an array, only add a new node to the top). Adding and removing elements just combines the array approach, until the array is full/empty, where it then performs the list operations. In hindsight, the list-array implementation is more like a compromise than the best of both worlds.

To get the top element we can store the total element count as length, then simply getting the head node gives us access to the array with the top element, so that lastly we




Figure 2.5: A simple graphical array example. The top element 'O' is also the last element in the array.

only need to do $index = (length - 1) \bmod chunk$, where *chunk* is the node's array size.

2.2 Array stack structure code

At the end I have decided to pick the array-based stack for one simple reason - continuous memory layout. I want to allow users to sort elements in the stack in order to have the same functionality as a priority queue, i.e. retrieving elements based on a priority, like the minimum (smallest) element. Of course the priority would be in reverse since the stack would get the last/top element in the array. It is possible to sort elements using the list and array-list implementations, but that requires creating a temporary array to copy elements back and forth.

The next chapters of this part will implement linked list restricted sequential structures. In chapter 4 the double ended queue data structure, which is implemented like a list of arrays, can also be used as a stack alternative with the list-array memory advantages.

2.2.1 The infinite stack structure

The `istack` code uses a dynamic array of elements based on a user-defined size for a singular element, for example using `sizeof` on `int`. The three members `size`, `length` and `capacity` represent single element size, current stack length and maximum capacity before resizing is required. The memory allocator allocates, reallocates and frees the elements array, and other temporary structures used in some functionalities.




Figure 2.6: Graphical example of singly-linked list of arrays. The top element is the last element in the head node's array.

```
include/sequence/istack.h
```

```
/// @brief Infinite stack data structure.
typedef struct infinite_stack {
    char * elements;           // array of elements
    size_t size, length, capacity; // size of single element,
    // structure length and its capacity
    memory_s const * allocator;
} istack_s;
```

2.2.2 Resize stack

Infinite expansion and shrinking require the understanding of reallocating memory and implementation of the `_istack_resize` function.

```
include/sequence/istack.c
```

```
/// @brief Resizes (reallocates) stack parameter arrays based on
    // changed capacity.
/// @param stack Structure to resize.
/// @param size New size.
void _istack_resize(istack_s * const stack, size_t const size);
```

The function calls on the allocator's `realloc` when high- and low-end capacity is reached, thus preserving minimum chunk length.

source/sequence/istack.c

```
{
    stack->capacity = size;

    stack->elements = stack->allocator->realloc(stack->elements, stack
->capacity * stack->size, stack->allocator->arg);
    // ...
}
```

2.2.3 Create stack

The `create_istack` function creates an empty infinite stack structure storing the element size and allocator.

include/sequence/istack.h

```
/// @brief Creates an empty structure.
/// @param size Size of a single element.
/// @return Stack structure.
istack_s create_istack(size_t const size);
```

The creation process uses a compound literal to quickly create and return a new infinite stack structure. It only needs to set the single element size and allocator members.

source/sequence/istack.c

```
{
    // ...
    return (istack_s) { .size = size, .allocator = &standard };
}
```

2.2.4 Make stack

The `make_istack` function, similar to `create_istack`, creates an infinite stack structure, but with the ability to specify a custom memory allocator.

include/sequence/istack.h

```
/// @brief Creates a custom empty structure.
/// @param size Size of a single element.
/// @param allocator Custom allocator structure.
/// @return Stack structure.
```

```
istack_s make_istack(size_t const size, memory_s const * const
    allocator);
```

The code is basically the same as the aforementioned function, but with standard allocator replaced by the parameter one.

source/sequence/istack.c

```
{
    // ...
    return (istack_s) { .size = size, .allocator = allocator };
}
```

2.2.5 Destroy stack

The `destroy_istack` destroys the stack and all its elements while at the same time making it invalid for future use. The parameters include a pointer to the stack, a function pointer taking in each element, and the function pointer's arguments, allowing the user to manipulate elements prior to their destruction. The destroy function pointer can be used to free memory, like strings, from the heap.

include/sequence/istack.h

```
/// @brief Destroys a structure and its elements, but makes it
    unusable.
/// @param stack Structure to destroy.
/// @param destroy Function pointer to destroy a single element.
/// @param ad Arguments for destroy function pointer.
void destroy_istack(istack_s * const stack, set_fn const destroy, void
    * const ad)
```

Since our elements are stored in an array, i.e. continuous memory, it is possible to write a for loop with a pointer representing a single element to call the destroy function pointer on. The variable then gets incremented by the specified size until the last valid memory address is reached.

source/sequence/istack.c

```
{
    // ...
    // for each element in elements array call destroy
    for (char * e = stack->elements; e < stack->elements + (stack->
length * stack->size); e += stack->size) {
        destroy(e, ad);
    }
}
```

```
// ...
```

After destroying each element the next step is to free the element's array using the specified allocator struct member.

source/sequence/istack.c

```
// ...
stack->allocator->free(stack->elements, stack->allocator->arg); //
free elements array
// ...
```

All that's left is to invalidate the stack, or set everything to zero, via calling `memset` from `<string.h>`.

source/sequence/istack.c

```
// ...
memset(stack, 0, sizeof(istack_s));
}
```

The `memset` function allows us to effectively set every member of a struct or any array to zero. It should only be used to set values to either all zero or all one (by using `-1` through `memset`'s second parameter).

Listing 2.1: The `memset` function prototype.

```
#include <string.h>

void *memset(void s[n], int c, size_t n);
```

Linux manual page - `memset(3)`

The `memset()` function fills the first `n` bytes of the memory area pointed to by `s` with the constant byte `c`.

2.2.6 Clear stack

Resembling `destroy_istack`, `clear_istack` can also be used to destroy each element in our structure, but the stack remains valid for future use.

include/sequence/istack.h

```
/// @brief Clears a structure and destroys its elements, but remains
usable.
/// @param stack Structure to destroy.
```

```

// @param destroy Function pointer to destroy a single element.
// @param ad Arguments for destroy function pointer.
void clear_istack(istack_s * const stack, set_fn const destroy, void *
    const ad);

```

This is achieved by freeing the elements array and setting the stack's length and capacity to zero, and elements array to NULL.

source/sequence/istack.c

```

{
    // for each element in elements array call destroy
    for (char * e = stack->elements; e < stack->elements + (stack->
length * stack->size); e += stack->size) {
        destroy(e, ad);
    }
    stack->allocator->free(stack->elements, stack->allocator->arg); //
    free elements array

    // set length to zero
    stack->length = stack->capacity = 0;
    stack->elements = NULL;
}

```

2.2.7 Copy stack

Creating a copy can be done with the `copy_istack` function. The function pointer's functionality is similar to the `memcpy` function in `<string.h>`, which takes a destination address, source address, and the size of the memory area to copy, see 2.2. However, the copy function pointer takes in `create_istack`'s last parameter as its third own parameter.

Listing 2.2: The `memcpy` function prototype from the Linux/UNIX manual page.

```
#include <string.h>
```

```

void *memcpy(void dest[restrict n], const void src[restrict n], size_t
    n);

```

Linux manual page - `memcpy(3)`

The `memcpy()` function copies `n` bytes from memory area `src` to memory area `dest`. The memory areas must not overlap. Use `memmove(3)` if the memory

areas do overlap.

include/sequence/istack.h

```

/// @brief Creates a copy of a structure and all its elements.
/// @param stack Structure to copy.
/// @param copy Function pointer to create a deep/shallow copy of a
///         single element.
/// @param ac Arguments for copy function pointer.
/// @return Stack structure.
istack_s copy_istack(istack_s const * const stack, copy_fn const copy,
    void * const ac);

```

First and foremost, a replica stack structure is created and initialized, with all members, except the elements array, being directly copied. Then the original allocator is used to generate a memory array.

source/sequence/istack.c

```

{
    // ...
    // create replica to initialize and return
    istack_s const replica = {
        .capacity = stack->capacity, .length = stack->length, .size =
        stack->size, .allocator = stack->allocator,
        .elements = stack->allocator->alloc(stack->capacity * stack->
        size, stack->allocator->arg),
    };
    // ...
}

```

All elements are later copied individually from source stack's memory at offset index position into replica's elements array.

source/sequence/istack.c

```

    // initialize replica's elements array with stack's elements
    for (size_t i = 0; i < stack->length; ++i) {
        size_t const offset = i * stack->size;
        copy(replica.elements + offset, stack->elements + offset, ac);
    }

    return replica;
}

```

2.2.8 Is empty stack

Since we can't remove elements from an empty stack, the function available to aid us in this endeavor is the `is_empty_istack`. It only returns a negated length.

include/sequence/istack.h

```
/// @brief Checks if structure is empty.
/// @param stack Structure to check.
/// @return 'true' if empty, 'false' if not.
bool is_empty_istack(istack_s const * const stack);
```

source/sequence/istack.c

```
{
    // ...
    return !(stack->length);
}
```

2.2.9 Push to stack

To add an element one can use the `push_istack` function. Since the stack is implemented using a resizable array checking if maximum capacity has been reached is a must. The index of where to push the element can be gained through the stack length member.

include/sequence/istack.h

```
/// @brief Pushes a single element to the top of the structure.
/// @param stack Structure to push into.
/// @param element Element buffer to push.
void push_istack(istack_s * const stack, void const * const element);
```

Firstly, we check if stack's length has reached the capacity for resizing into a new chunk of memory. Since the structures grow exponentially, starting at a specific power of two chunk, it is important to resize it into `ISTACK_CHUNK` if length member is zero or else continue with doubling.

source/sequence/istack.c

```
{
    // ...
    if (stack->length == stack->capacity) { // if length is equal to
        capacity the array must expand linearly
        size_t const capacity = stack->length ? stack->length *
        CERPEC_FACTOR : ISTACK_CHUNK;
```

```

        _istack_resize(stack, capacity);
    }
    // ...

```

As the size of a single element is user-define, the only optimal way to copy the element into our structure is to use `memcpy`. Because the `element` parameter and `elements` stack member are stored in two different memory locations `memcpy` and its speed can be utilized.

source/sequence/istack.c

```

    // ...
    // push element knowing the elements array can fit it
    memcpy(stack->elements + (stack->length * stack->size), element,
stack->size);
    stack->length++;
}

```

2.2.10 Pop from stack

Removal of elements is done through the `pop_istack` function. If the structure is empty the program logically errors.

include/sequence/istack.h

```

/// @brief Pops a single element from the top of the structure.
/// @param stack Structure to pop from.
/// @param buffer Element buffer to save pop.
void pop_istack(istack_s * const stack, void * const buffer);

```

Popping the top element is as simple as getting the element at decremented length, using it as the index position and copying into a temporary variable parameter.

source/sequence/istack.c

```

{
    // ...
    // remove element from elements array
    stack->length--;
    memcpy(buffer, stack->elements + (stack->length * stack->size),
stack->size);
}

```

The structure adjusts itself during growth, therefore it is important to not waste memory. After reaching a smaller power of two capacity the stack automatically shrinks. Chunk length is a special case where if the capacity goes below it, the memory array stays the same, but only until reaching zero. After that it is set to zero.

source/sequence/istack.c

```
// ...
if (stack->length ≤ stack->capacity / CERPEC_FACTOR && (stack->
length > ISTACK_CHUNK || !stack->length)) {
    _istack_resize(stack, stack->length);
}
}
```

2.2.11 Peep the stack

Peeping the last added element via `peep_istack` is a little bit like popping it.

include/sequence/istack.h

```
/// @brief Peeps a single element from the top of the structure.
/// @param stack Structure to peep.
/// @param buffer Element buffer to save peep.
void peep_istack(istack_s const * const stack, void * const buffer);
```

The idea is to just get the element index at $length - 1$ and save it into a temporary parameter.

source/sequence/istack.c

```
{
    // ...
    // only copy the top element into the buffer
    memcpy(buffer, stack->elements + ((stack->length - 1) * stack->
size), stack->size);
}
```

2.2.12 For each element

Each element can be accessed and manipulated through the special `each_istack` function. This function iterates not at the top, but at the bottom, while making its way up to the top. The reason it starts at the bottom is because it makes the code iterating through the values shorter.

include/sequence/istack.h

```
/// @brief Iterates over each element in structure starting from the
beginning.
/// @param stack Structure to iterate over.
```

```

/// @param manage Function pointer to manage each element reference
/// using generic arguments.
/// @param am Generic arguments to use in function pointer.
void each_istack(istack_s const * const stack, manage_fn const manage,
    void * const am);

```

All the elements are in an array, thus iterating over them is as simple as using an empty for loop one-liner. The manage function pointer, together with its arguments, allows direct access to the element, making manipulation is possible. If manage returns false, the iteration automatically stops, kinda like the break keyword in a classic for loop.

source/sequence/istack.c

```

{
    // ...
    // iterate over each element from bottom to the top of stack
    for (char * e = stack->elements; e < stack->elements + (stack->
length * stack->size) && manage(e, am); e += stack->size) {}
}

```

2.2.13 Apply to all elements

The apply_istack gives access to all the elements as if they were inside a continuous array, which is really simple for this stacks, but it is extremely hard to implement for structures in the next chapters - Queue and Deque. The main use of apply_istack is to sort the array of elements as fast as the sorting function pointer process and its generic arguments allow.

include/sequence/istack.h

```

/// @brief Apply each element in structure into an array to manage.
/// @param stack Structure to map.
/// @param process Function pointer to process array of elements using
/// structure length and arguments.
/// @param ap Generic arguments to use in function pointer.
void apply_istack(istack_s const * const stack, process_fn const
    process, void * const ap);

```

Similarly to the each_istack, the apply_istack only has one line of code - the process function pointer itself.

source/sequence/istack.c

```

{

```

```
// ...  
// process stack elements as an array (as a whole)  
process(stack->elements, stack->length, ap);  
}
```

2.3 Example usages

Here are some example usages of the infinite stack. Each example show a stack of integers performing some of the functionalities mentioned in the previous section.

2.3.1 Creating clearing and destroying a stack

This example shows how stack creation, clearing and destroying are performed in code.

```
#include <sequence/istack.h>  
  
void intdst(void * const element, void * null) {  
    (void)(null);  
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);  
}  
  
int main(void) {  
    istack_s stack = create_istack(sizeof(int));  
  
    // some operations  
  
    clear_istack(&stack, intdst, NULL);  
  
    // some new operations  
  
    destroy_istack(&stack, intdst, NULL);  
  
    return 0;  
}
```

2.3.2 Pushing, peeping and popping elements

Code showing push, peep and pop example for 10 integer elements.

```
#include <stdio.h>  
#include <sequence/istack.h>
```

```

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

int main(void) {
    istack_s stack = create_istack(sizeof(int));

    // push 10 elements
    for(int i = 0; i < 10; i++) {
        push_istack(&stack, &i);
    }

    // see if top element is 9
    int v = -1;
    peep_istack(&stack, &v);
    printf("%d\n", v);

    // pop and print all elements
    while(!is_empty_istack(&stack)) {
        pop_istack(&stack, &v);
        printf("%d ", v);
    }

    destroy_istack(&stack, intdst, NULL);

    return 0;
}

```

2.3.3 Iterating and sorting elements

Iterating and applying a sorting function to the stack which prints each element, then reverses the ordering in the structure, and lastly prints the stack again, but with elements in reverse order.

```

#include <stdio.h>
#include <stdlib.h>
#include <sequence/istack.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

```

```
}

int intrcmp(void * const a, void * const b) {
    return (*(int*)b) - (*(int*)a); // reverse comparison
}

bool intprt(void * const element, void * format) {
    printf(format, (*(int*)element));
    return true;
}

void intsrt(void * const array, size_t const length, void * const
    compare) {
    qsort(array, length, sizeof(int), compare);
}

int main(void) {
    istack_s stack = create_istack(sizeof(int));

    // push 10 elements
    for(int i = 0; i < 10; i++) {
        push_istack(&stack, &i);
    }

    // print each element from bottom to top in order
    each_istack(&stack, intprt, "%d ");

    // sort elements in reverse
    // YOU SHOULD NOT CAST A FUNCTION POINTERS INTO VOID*
    // ARGUMENTS, THIS IS JUST A SILLY EXAMPLE, INSTEAD PUT
    // THE FUNCTION POINTER INSIDE A STRUCT AND PUT THE
    // INITIALIZED STRUCT POINTER ITSELF AS AN ARGUMENT
    apply_istack(&stack, intsrt, intrcmp);
    puts("");

    // print each element from bottom to top in reverse
    each_istack(&stack, intprt, "%d ");

    destroy_istack(&stack, intdst, NULL);

    return 0;
}
```

2.4 On a final note...

Starting simply, this stack implementation internet article [2] is a good entry source explaining how stacks work in a way shorter format compared to this book. [1] is the C++ standard library implementation of the stack. this implementation is merely an inherited template class of the deque data structure, which thankfully proves that this book's deque can also be used as a stack.

Further reading

- [1] cppreference.com. std::stack. <https://en.cppreference.com/cpp/container/stack>. Revision ID: 183594. (Visited on 07/09/2026).
- [2] GeeksforGeeks. Stack Data Structure. <https://www.geeksforgeeks.org/dsa/stack-data-structure/>. Last updated: 20 Jan, 2026. Jan. 2026. (Visited on 07/09/2026).

Chapter 3

Queue

Imagine you're waiting for your daily calorie intake, i.e. lunch, in school and you see a queue. When one person gets their lunch the entire queue moves one person to the front, because logically, the entire cafeteria won't move to the next person.

But that's not how queues in programming work, I mean, technically they can, but//... If each person is a single element and the cafeteria is at the current place of the first person in the queue, then you don't want to move each element to the cafeteria, but move the cafeteria itself to the next person. So we have a current pointer/index/position which, after getting rid of the first element, we go/increment/move to the second one, and so on.

Going back to chapter 2's chess analogy, before we put those moves into the stack we're kind of pulling them from an array of moves starting at the front until the back is reached. Now lets say there're more moves still happening, plus we want to play back and forth the moves which were already played without waiting for the game to finish. To solve this we can pipe the moves into a special queue as they happen, then push elements from the queue into the first stack. If we want to go to the beginning, then we push the moves from the first stack into the second. To play all the way to the latest move, all that's left is to first push all the elements from the second into the first stack, and push all moves from the queue into first. The current saved element is the latest move, see Figure 3.1.

3.1 Implementations of a queue

Similarly to stacks, there are also ~~three~~ two ways to implement queues:

- Singly linked lists
- Static arrays (not dynamic)




Figure 3.1: Visual example of two stacks and a queue being used to track the chess game.




Figure 3.2: Visual representation of a circular singly linked list.

3.1.1 The free linked list

Usually when implementing a queue using linked lists we create a list structure that has one reference to the head - for access and removal, and to the tail node - for pushing elements.

However, there is also a slightly better way to implement them, which is to use circular lists. The next subsection talks about circular arrays, but we can also use circular lists, kinda like in Figure 3.2. Hear me out, this won't work if we only have a reference to the head, so let's go with the next best thing - the tail. Since the list is circular, tail's next node will therefore be the head, or front node.

This makes access and removal as simple as getting the head from the tail, also adding elements is easy since we just make the new node's next reference the head and change tail's next to new node, just see Figure 3.3. The downside is that circular linked lists are more complex than the classical list implementation, but by just having one tail node pointer we can decrease the size of the structure by quite a bit (or a whopping 32/64 bits at that).

3.1.2 The confined array

Another way to implement a queue is to use a static array. We can use dynamic arrays as well, but when the queue needs to expand the elements must be copied front to back into a new array. This sounds reasonable for array resizing, but usually when arrays resize they don't always need to copy all the elements and just either don't do anything or they just inform that the memory length was changed. This is the case for chapter 2's stacks, but in the case of queue a copy needs to happen all the time to move the elements to the beginning of the array, which is a huge performance hit. But what if we add the simplicity of a static array queue with the expandability of tail-based circular linked lists?




Figure 3.3: Example of adding (top) and removing (bottom) an element from a circular singly linked list queue.




Figure 3.4: Graphical example of tail-based circular linked list of arrays. The front/fore element is the first element (index one) in the head node's array, and the back element is the last element (index six) in the tail node's array.

3.1.3 Tail-based circular linked list of arrays

So we have a tail-based circular linked list with nodes, but instead of the node containing only one element it has an entire array of elements. Technically, a simple linked list with one node element is a subset of all possible tail-based circular linked lists. So we have an array of continuous memory for each element, we get rid of the resizing and recopying from the start problem, and last, but not least, kinda/... we may get better performance for big sizes since it's simpler to allocate a new node of size

$$136b = (8b + 128b) = (\text{sizeof}(\text{next_pointer}) + (32 * \text{sizeof}(\text{integer_element})))$$

than it is to resize an elements array from 4MiB to 8MiB with potential copying. The downside is that the code complexity of tail-based circular linked lists and circular static arrays add up when combined together.

3.2 Tail-based circular linked list of array code

Since I promised in chapter 2 to implement a linked list implementation of a restricted sequential data structure I've decided to do it with the queue (spoilers if you didn't read chapter 2: and also the Deque).

3.2.1 The infinite queue structure

The implementation is made up of two structures, let's call them the node and the queue.

The node uses a less known C feature - flexible array members, which are arrays without a specified size at the end of a structure - the structure is thus pointer to next node and empty array member.

The queue structure has a tail node pointer, the classic element size and structure length, and an allocator reference. It also includes a current index which we haven't talked about. The current index represents the position of the front element in head relative to the array, so when the index reaches the end it is set to zero. To get the back element we just add the current index with the queue length and multiply the result with the element size.

GNU C Language Manual - Flexible Array Fields

The C99 standard adopted a more complex equivalent of zero-length array fields. It's called a flexible array, and it's indicated by omitting the length, like this:

```
struct line { int length; char contents[]; };
```

The flexible array has to be the last field in the structure, and there must be other fields before it.

Under the C standard, a structure with a flexible array can't be part of another structure, and can't be an element of an array.

GNU C allows static initialization of flexible array fields. The effect is to "make the array long enough" for the initializer.

```
struct f1 { int x; int y[]; } f1 = { 1, { 2, 3, 4 } };
```

This defines a structure variable named f1 whose type is struct f1. In C, a variable name or function name never conflicts with a structure type tag.

Omitting the flexible array field's size lets the initializer determine it. This is allowed only when the flexible array is defined in the outermost structure and you declare a variable of that structure type. For example:

```
struct foo { int x; int y[]; }; struct bar { struct foo z; };
```

```
struct foo a = { 1, { 2, 3, 4 } }; // Valid.  
struct bar b = { { 1, { 2, 3, 4 } } }; // Invalid.  
struct bar c = { { 1, { } } }; // Valid.
```

```
struct foo d[1] = { { 1 { 2, 3, 4 } } }; // Invalid.
```

```
include/sequence/iqueue.h
```

```
/// @brief Circular linked list node for queue structure.
struct infinite_queue_node {
    struct infinite_queue_node * next; // next sibling node
    char elements[];                  // flexible elements array with
    IQUEUE_CHUNK length
};

/// @brief Infinite queue data structure.
typedef struct infinite_queue {
    struct infinite_queue_node * tail; // tail node to append next
    elements while enqueue-ing
    size_t size, current, length;      // current index, element size
    and structure length
    memory_s const * allocator;
} iqueue_s;
```

3.2.2 Create queue

The *create_iqueue* function create an empty circular list queue structure storing the allocator, index and size data.

```
include/sequence/iqueue.h
```

```
/// @brief Creates an empty structure.
/// @param size Size of a single element
/// @return Queue structure.
iqueue_s create_iqueue(size_t const size);
```

The creation process uses a compound literal to quickly create and return a new infinite queue structure. It only needs to set the single element size and allocator members.

```
source/sequence/iqueue.c
```

```
{
    // ...
    return (iqueue_s) { .size = size, .allocator = &standard };
}
```

3.2.3 Make queue

The `make_iqueue` function, similar to `create_iqueue`, creates an infinite queue structure, but with the ability to specify a custom memory allocator.

include/sequence/iqueue.h

```
/// @brief Creates a custom empty structure.
/// @param size Size of a single element.
/// @param allocator Custom allocator structure.
/// @return Queue structure.
iqueue_s make_iqueue(size_t const size, memory_s const * const
    allocator);
```

The code is basically the same as the aforementioned function, but with the standard allocator replaced with the parameter one.

source/sequence/iqueue.c

```
{
    // ...
    return (iqueue_s) { .size = size, .allocator = allocator };
}
```

3.2.4 Destroy queue

Next up, the `destroy_iqueue` function destroys the queue structure and all of its nodes and elements. The element destruction is done by the `set_fn destroy` function pointer, which is called for each element before the node is freed. The node destruction is done by iterating through the circular list - not from the tail, but the head; and freeing each node until we get back to the tail. The code is a little bit more complex because of circularity and array node members.

include/sequence/iqueue.h

```
/// @brief Destroys a structure and its elements, but makes it
    unusable.
/// @param queue Structure to destroy.
/// @param destroy Function pointer to destroy a single element.
/// @param ad Arguments for destroy function pointer.
void destroy_iqueue(iqueue_s * const queue, set_fn const destroy, void
    * const ad);
```

Since our elements are stored in a linked list it is possible to iterate via going to the next node. The head node is the next node from the tail, so in order to begin destruction at the front we must reference the current node from the previous (tail) one.

source/sequence/iqueue.c

```
{
    // ...
    // for each node starting from tail (previous to head)
    struct infinite_queue_node * previous = queue->tail;
    for (size_t start = queue->current; queue->length; start = 0) {
        // ...
```

After that the array loop follows, which has a more complex logic, since the head and tail node can have different numbers of elements, unlike the in-between nodes which are always full. There can also be the case that only one node contains all elements, i.e. head equals tail, and it may not be full and the first and last few element slots may be empty.

source/sequence/iqueue.c

```
    // ...
    // destroy elements in previous' next (so head and next)
    size_t i = start;
    for (; i < queue->length && i < IQUEUE_CHUNK; ++i) {
        destroy(previous->next->elements + (i * queue->size), ad);
    }
    queue->length -= (i - start);
    // ...
```

All that's left in the loop is to free the node and move to the next one until we get back to the tail, which is the last node to be freed.

source/sequence/iqueue.c

```
    // save destroyed node and cut it from queue list
    struct infinite_queue_node * temp = previous->next;
    previous->next = previous->next->next;

    // free destroyed node
    queue->allocator->free(temp, queue->allocator->arguments);
}
```

Lastly we use `memset` to invalidate the queue from future use.

source/sequence/iqueue.c

```
    // set everything to zero
    memset(queue, 0, sizeof(iqueue_s));
}
```

3.2.5 Clear queue

Similar to the destruction function, the `clear_iqueue` function clears the queue of all elements and nodes, but keeps the structure itself usable. The code is basically the same as the destruction function, but without the final `memset` and with only some members set to zero/NULL.

include/sequence/iqueue.h

```

/// @brief Clears a structure and destroys its elements, but remains
        usable.
/// @param queue Structure to destroy.
/// @param destroy Function pointer to destroy a single element.
/// @param ad Arguments for destroy function pointer.
void clear_iqueue(iqueue_s * const queue, set_fn const destroy, void *
        const ad);

```

source/sequence/iqueue.c

```

{
    // ...
    // for each node starting from tail (previous to head)
    struct infinite_queue_node * previous = queue->tail;
    for (size_t start = queue->current; queue->length; start = 0) {
        // destroy elements in previous' next (so head and next)
        size_t i = start;
        for (; i < queue->length && i < IQUEUE_CHUNK; ++i) {
            destroy(previous->next->elements + (i * queue->size), ad);
        }
        queue->length -= (i - start);

        // save destroyed node and cut it from queue list
        struct infinite_queue_node * temp = previous->next;
        previous->next = previous->next->next;

        // free destroyed node
        queue->allocator->free(temp, queue->allocator->arguments);
    }

    queue->current = queue->length = 0;
    queue->tail = NULL;
}

```

3.2.6 Copy queue

Continuing with the high complexity, the next function is the `copy_iqueue` function, which creates a copy of the queue structure and all of its nodes and elements together with user defined arguments. The element copying is done by the `copy_fn` copy function pointer, which is called for each element to create either a deep or shallow copy.

include/sequence/iqueue.h

```

/// @brief Creates a copy of a structure and all its elements.
/// @param queue Structure to copy.
/// @param copy Function pointer to create a deep/shallow copy of a
    single element.
/// @param ac Arguments for copy function pointer.
/// @return Queue structure.
iqueue_s copy_iqueue(iqueue_s const * const queue, copy_fn const copy,
    void * const ac);

```

First we initialize an empty queue replica.

source/sequence/iqueue.c

```

{
    // ...
    // create properly initialized replica
    iqueue_s replica = {
        .size = queue->size, .length = queue->length, .current = queue
->current, .allocator = queue->allocator,
    };
    // ...

```

Then we create a current pointer to queue's tail and a double pointer to replica's tail so as to be able to change replica's tail and its reference, and remove the special head case.

source/sequence/iqueue.c

```

    // ...
    // set original queue's and replica's current nodes for iteration
    struct infinite_queue_node const * current_queue = queue->tail;
    struct infinite_queue_node ** current_copy = &(replica.tail); //
two pointer list to remove special head case
    // ...

```

The loop starts with initializing the remaining elements count and starting index, like in `destroy_iqueue` and `clear_iqueue`. The loop ends when there are no more remaining elements to copy. The last pointer just contains the last created node to make

the list circular and redirect the replica's tail to it since it would contain the head after the loop ends.

source/sequence/iqueue.c

```
// ...
struct infinite_queue_node * last = NULL;
for (size_t remaining = queue->length, start = queue->current;
remaining; start = 0) {
// ...
```

The first part of the loop creates a new empty node and redirects replica's and last's node to it.

source/sequence/iqueue.c

```
// ...
// allocate new node for replica
last = (*current_copy) = queue->allocator->alloc(sizeof(struct
infinite_queue_node) +
(IQUEUE_CHUNK * queue->size), queue->allocator->arg);
// ...
```

The second part just makes more redirects for replica's node and changes the queue's current node to the next.

source/sequence/iqueue.c

```
// ...
// redirect current replica's node with replica's tail for
circularity
(*current_copy)->next = replica.tail;
current_queue = current_queue->next;
// ...
```

The third part iterates over the node array and creates copies based on the copy function pointer. The logic is more complex since elements may not lay at the beginnings/ends of the array.

source/sequence/iqueue.c

```
// ...
// outside for loop to get copied chunk size, since it can
either be 'remaining' or 'IQUEUE_CHUNK'
size_t i = start;
for (; i < remaining && i < IQUEUE_CHUNK; ++i) { // while i is
less than both 'remaining' and 'IQUEUE_CHUNK'
```

```

        copy((*current_copy)->elements + (i * queue->size),
current_queue->elements + (i * queue->size), ac);
    }
    remaining -= i; // subtract copied size from remaining size
using i
    // ...

```

Last part of the loop only redirects to the next replica node.

source/sequence/iqueue.c

```

    // ...
    current_copy = &((*current_copy)->next);
}
// ...

```

And with that the loop ends and all that's left is to properly redirect the replica's tail to the last node and return the initialized structure copy.

source/sequence/iqueue.c

```

    // ...
    replica.tail = last; // set replica's tail to its last added node
    (i.e. its tail)

    return replica;
}

```

3.2.7 Is queue empty

Yep, the `is_empty_iqueue` function checks if the queue is empty by checking if the length is zero.

include/sequence/iqueue.h

```

/// @brief Checks if structure is empty.
/// @param queue Structure to check.
/// @return 'true' if empty, 'false' if not.
bool is_empty_iqueue(iqueue_s const * const queue);

```

source/sequence/iqueue.c

```

{
    // ...
    return (queue->length == 0);
}

```

3.2.8 Enqueue

The `enqueue_iqueue` function adds an element to the back of the queue. The code is a little bit more complex because of the circular list and array members, but the logic is pretty simple - we check if the tail node is full, if it is we create a new one and redirect the tail to it, then we add the element to the back of the tail node's array and increment the length.

include/sequence/iqueue.h

```
/// @brief Enqueues a single element to the end of the structure.
/// @param queue Structure to enqueue into.
/// @param element Element buffer to enqueue.
void enqueue_iqueue(iqueue_s * const queue, void const * const element
);
```

First we check if head's array is full, if yes then we must create a new node and perform pointer redirects, else we just copy the element to the node array.

source/sequence/iqueue.c

```
{
    // ...
    // index where the next element will be enqueued
    size_t const next_index = (queue->current + queue->length) %
    IQUEUE_CHUNK;
    if (!next_index) { // if head list array is full (is divisible)
        adds new list element to head
    }
    // ...
```

Inside the if statement we create a new node.

source/sequence/iqueue.c

```
    // ...
    struct infinite_queue_node * node = queue->allocator->alloc(
    sizeof(struct infinite_queue_node) + (IQUEUE_CHUNK * queue->size),
    queue->allocator->arguments);
    // ...
```

At the end of the condition we redirect the newly created node and queue's tail properly.

source/sequence/iqueue.c

```
    // ...
    if (queue->tail == NULL) {
        node->next = node; // create initial circle
```

```

    } else {
        node->next = queue->tail->next; // make temp's next node
head node
        queue->tail->next = node; // make previous tail's next
node point to temp
    }
    queue->tail = node;
}
// ...

```

At last, we copy the element to its proper place in its proper node and increment the length.

source/sequence/iqueue.c

```

// ...
memcpy(queue->tail->elements + (next_index * queue->size), element,
queue->size);
queue->length++;
}

```

3.2.9 Dequeue

The `dequeue_iqueue` function removes the first added element from the queue's list node at head.

include/sequence/iqueue.h

```

/// @brief Dequeues a single element from the start of the structure.
/// @param queue Structure to dequeue from.
/// @param buffer Element buffer to save dequeue.
void dequeue_iqueue(iqueue_s * const queue, void * const buffer);

```

Firstly the head node is reached and its first element is copied to the buffer, then queue's length is decremented and new current index is calculated.

source/sequence/iqueue.c

```

{
    // ...
    memcpy(buffer, queue->tail->next->elements + (queue->current *
queue->size), queue->size);
    queue->length--; // decrement queue size
    queue->current = (queue->current + 1) % IQUEUE_CHUNK; // set
current to next index in node array
    // ...
}

```

If there are no more elements in the queue the last head/tail node is freed and everything is set to zero/NULL. Else if current has circled back to zero the current head node is empty and needs to be freed and removed. After that head's next node becomes the new head.

source/sequence/iqueue.c

```
// ...
if (!queue->length) { // if queue is empty after extracting
    element, then free memory and reset everything to zero
    queue->allocator->free(queue->tail, queue->allocator->
arguments); // free empty tail/head node

    queue->current = 0; // if queue is empty make current index 0
to not break enqueue_iqueue operation
    queue->tail = NULL; // set tail to NULL
} else if (queue->current == 0) { // else if current index circles
back, free start list element and shift to next
    struct infinite_queue_node * head = queue->tail->next; // get
empty head node
    queue->tail->next = queue->tail->next->next; // set new head
node to its next node

    queue->allocator->free(head, queue->allocator->arguments); //
free previous head node
}
}
```

3.2.10 Peek queue

peek_iqueue peek the front element at head and saves it into buffer without removing it from the queue.

include/sequence/iqueue.h

```
/// @brief Peeks a single element from the start of the structure.
/// @param queue Structure to peek.
/// @param buffer Element buffer to save peek.
void peek_iqueue(iqueue_s const * const queue, void * const buffer);
```

We just copy the current element at head into the buffer and return.

source/sequence/iqueue.c

```
{
```

```

    // ...
    memcpy(buffer, queue->tail->next->elements + (queue->current *
queue->size), queue->size);
}

```

3.2.11 For each element

Each element can be accessed and manipulated through the special `each_iqueue` function which iterates from the front to the back, making its way backwards. That is standard queue traversal behavior.

include/sequence/iqueue.h

```

/// @brief Iterates over each element in structure starting from the
beginning.
/// @param queue Structure to iterate over.
/// @param manage Function pointer to operate on each element
reference using element size and generic arguments.
/// @param am Generic arguments to use in function pointer.
void each_iqueue(iqueue_s const * const queue, manage_fn const manage,
void * const am);

```

Initially we need to save the remaining size as the queue's length and the previous node as the tail to start iteration from the head. Since head and tail may not have elements saved at the beginning/end of their arrays, we need to create the remaining size.

source/sequence/iqueue.c

```

{
    // ...
    size_t remaining = queue->length; // save queue size as remaining
size for iteration
    struct infinite_queue_node const * previous = queue->tail;

```

The for loop iterates over each node and element until remaining size is zero. Since the first inserted element is stored in head the current node is inferred from the previous node's next reference. The secondary for loop just makes sure to begin iteration from the current index position in array and then continue from zero while elements are managed. The manage function is checked and ends the main `each_iqueue` function when manage returns false.

source/sequence/iqueue.c

```

    // ...
    // while remaining size wasn't decremented to zero

```

```

for (size_t start = queue->current; remaining; start = 0, previous
    = previous->next) {
    struct infinite_queue_node * current = previous->next;

    size_t i = start; // save i outside loop to later use it in
    subtraction
    for (;i < remaining && i < IQUEUE_CHUNK; ++i) { // while i is
    less than either remaining size or node's array size
        // manage on element and if zero is returned then end main
        function
        if (!manage(current->elements + (i * queue->size), am)) {
            return;
        }
    }
    remaining -= (i - start); // subtract absolute value of i and
    start from remaining size
}
}

```

3.2.12 Apply to all elements

apply_iqueue applies a process function pointer on an array of elements, thus allowing sorting of queue elements. It is way more complex than the apply_istack function from the previous chapter.

include/sequence/iqueue.h

```

// @brief Apply each element in structure into an array to manage.
// @param queue Structure to map.
// @param process Function pointer to manage array of elements using
// structure length, element size and arguments.
// @param ap Generic arguments to use in function pointer.
void apply_iqueue(iqueue_s const * const queue, process_fn const
    process, void * const ap);

```

At the beginning we allocate a temporary array to save the elements from the queue's nodes. The array size is the queue's length multiplied by the element size.

source/sequence/iqueue.c

```

{
    // ...
    // create elements array to temporary save element from circular
    linked list nodes

```

```
char * elements_array = queue->allocator->alloc(queue->length *
queue->size, queue->allocator->arguments);
```

Then we have our first evil for loop, which iterates over the nodes and elements in the same way as `each_iqueue`, but instead of handling each element it saves them into the temporary array. The logic is complex.

source/sequence/iqueue.c

```
// ...
size_t index = 0, remaining = queue->length; // temporary variable
to save queue's size
struct infinite_queue_node const * previous = queue->tail; //
create node pointer to save previous nodes into
for (size_t start = queue->current; remaining; start = 0, previous
= previous->next) {
    size_t i = start; // extract i from for loop to later subtract
it from remaining size
    for (; i < remaining && i < IQUEUE_CHUNK; ++i) { // while i is
less than remaining size and node array size
        // save previous' next elements (so current) to elements
array at index, and increment index
        memcpy(elements_array + (index * queue->size), previous->
next->elements + (i * queue->size), queue->size);
        index++;
    }
    remaining -= (i - start); // subtract absolute value of i and
start from remaining size
}
// ...
```

Now we can finally process the elements as a single array.

source/sequence/iqueue.c

```
// ...
process(elements_array, queue->length, ap); // process initialized
elements array
// ...
```

Now we go to the second evil for loop, which basically just copies the elements from the array into the queue's nodes, thus saving the changes made by the process function. The logic is basically the same as the first loop, but instead of saving elements into the array it saves them back into the nodes.

source/sequence/iqueue.c

```

// ...
index = 0, remaining = queue->length, previous = queue->tail;
for (size_t start = queue->current; remaining; start = 0, previous
    = previous->next) { // save elements array back into queue
    size_t i = start; // extract i from for loop to later subtract
    it from remaining size
    for (; i < remaining && i < IQUEUE_CHUNK; ++i) { // while i is
    less than remaining size and node array size
        memcpy(previous->next->elements + (i * queue->size),
        elements_array + (index * queue->size), queue->size);
        index++;
    }
    remaining -= (i - start); // subtract absolute value of i and
    start from remaining size
}
// ...

```

And all that's left is to free the temporary array and return.

source/sequence/iqueue.c

```

// ...
queue->allocator->free(elements_array, queue->allocator->arguments
);
}

```

3.3 Example usages

Here are some example usages of the infinite queue. Each example show a queue of integers performing some of the functionalities mentioned in the previous section.

3.3.1 Creating, clearing and destroying a queue

Creation and clearing/destruction of a queue in code form. This example is more of a empty skeleton showcase of how one might use the three basic queue functions.

```

#include <sequence/iqueue.h>

```

```

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

```

```
int main(void) {
    iqueue_s queue = create_iqueue(sizeof(int));

    // some operations

    clear_iqueue(&queue, intdst, NULL);

    // some new operations

    destroy_iqueue(&queue, intdst, NULL);

    return 0;
}
```

3.3.2 Enqueueing, peeking and dequeuing elements

Example enqueueing, peeking and dequeuing integer elements from a queue. It shows a short code example of a useless usage for the queue library.

```
#include <stdio.h>
#include <sequence/iqueue.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

int main(void) {
    iqueue_s queue = create_iqueue(sizeof(int));

    // enqueue 10 elements
    for(int i = 0; i < 10; i++) {
        enqueue_iqueue(&queue, &i);
    }

    // see if front element is 0
    int v = -1;
    peek_iqueue(&queue, &v);
    printf("%d\n", v);

    // dequeue and print all elements
    while(!is_empty_iqueue(&queue)) {
```

```

        dequeue_iqueue(&queue, &v);
        printf("%d ", v);
    }

    destroy_iqueue(&queue, intdst, NULL);

    return 0;
}

```

3.3.3 Iterating and sorting elements

Iterate-prints through the inserted queue elements, sorts them based on `intsrt` function, and then iterate-prints the sorted elements, in reverse order.

```

#include <stdio.h>
#include <stdlib.h>
#include <sequence/iqueue.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

int intrcmp(void * const a, void * const b) {
    return (*(int*)b) - (*(int*)a); // reverse comparison
}

bool intprt(void * const element, void * format) {
    printf(format, (*(int*)element));
    return true;
}

void intsrt(void * const array, size_t const length, void * const
compare) {
    qsort(array, length, sizeof(int), compare);
}

int main(void) {
    iqueue_s queue = create_iqueue(sizeof(int));

    // enqueue 10 elements
    for(int i = 0; i < 10; i++) {
        enqueue_iqueue(&queue, &i);
    }
}

```

```

}

// print each element from start to end in order
each_iqueue(&queue, intprt, "%d ");

// sort elements in reverse
// YOU SHOULD NOT CAST A FUNCTION POINTERS INTO VOID*
// ARGUMENTS, THIS IS JUST A SILLY EXAMPLE, INSTEAD PUT
// THE FUNCTION POINTER INSIDE A STRUCT AND PUT THE
// INITIALIZED STRUCT POINTER ITSELF AS AN ARGUMENT
apply_iqueue(&queue, intsrt, intrcmp);
puts("");

// print each element from start to end in reverse
each_iqueue(&queue, intprt, "%d ");

destroy_iqueue(&queue, intdst, NULL);

return 0;
}

```

3.4 On a final note...

The [2] article again simply explains what queue's are. C++'s [1] shows how queues are also just dequeues in a sheep's clothing.

Further reading

- [1] cppreference.com. std::queue. <https://en.cppreference.com/cpp/container/queue>. [Online; accessed 20-April-2026]. 2026.
- [2] GeeksforGeeks. Queue Data Structure. <https://www.geeksforgeeks.org/dsa/queue-data-structure/>. Last updated: 20 Jan, 2026. Jan. 2026. (Visited on 07/09/2026).

Chapter 4

Deque

By imagining a music player, it has a list of songs which the user can play. To play the songs a user needs to interact with a music player interface. The user can click play, but when they don't like the initial song they can also interact with a next (or previous) button to change songs. When the next song is selected, what should happen when the previous button is clicked? Of course, the previously loaded song should start playing. To implement this functionality a programmer can employ a deque (pronounced /dek/ DEK).

Deques allow data insertion and removal from both ends, which can be a little bit confusing to wrap your head around, but that is the functionality needed for such projects.

Going back to chapter 3's playback chess player, since the game is short, it will be quite fast to push all moves to the queue. Thus it may be more convenient to use a deque for smaller sets of moves. We have a deque with all chess moves inside it, so to play back and forth all that is needed is to have a full deque and an integer variable that stores the number of steps taken and a sign representing the direction from the initial/first move (since players don't usually allow circling to the other end when the current pointer lies on the edge). You can see Figure 4.1 for a better understanding.

4.1 Implementations of a deque

The deque also has three two implementations. Those implementations are similar to the queue's, but in the case of the array we access elements from two ends. Meanwhile the list implementation relies on doubly linked lists for both two-end access and traversal:

- Doubly-linked lists
- Static arrays (not dynamic)




Figure 4.1: Visual example of a deque with chess moves and a current index allowing element traversal from first-to-last (top two) and last-to-first (bottom two).

4.1.1 The simple array

The simplest implementation can include an array of objects and an index pointing to the current objects. To move through them, all that's needed is to increment/decrement the index. But what if the index is zero and decrementing causes the index to become negative, or even worse - underflow? The answer is simple, check if index is zero and then circle back, like this $current = length - 1$. The simple array-based deque also needs to store the object count - deque length. Overflowing is fixed by checking if the current index is less than deque's length and doing the reverse, by making $current = 0$. Visual example shown in Figure 4.2.




Figure 4.2: Visual example of an array-based deque with a current front (fore) index allowing element traversal from first-to-last (fore-to-back) and last-to-first (back-to-fore).




Figure 4.3: Visual example of a list-based deque with a current front (fore) head node allowing element traversal from first-to-last (fore-to-back) and last-to-first (back-to-fore).

4.1.2 The complex doubly-linked list

The most complex implementation can include linked list with one previous and one next pointer. To move through nodes simply get either the head (front) or tail (back) node and iterate through the entire list. To save memory we can make the list circular and only store the head node, since the tail will be accessible through head's previous node pointer, see Figure 4.3.

Another memory optimization for doubly-linked lists are XOR linked lists, which store the next and previous pointers as a single XOR pointer of the two, or more accurately, $combined(A) = next(A) \oplus previous(A)$. To retrieve the next (or previous) node pointer, simply save the head (and tail) pointers as they are and XOR the head's (or tail's) combined pointer, like this $next(A) = combined(A) \oplus previous(A)$. Again, to save memory we add more operations, but a drawback is that the list is harder to debug. C doesn't allow bitwise operations on pointers, but this can be avoided by casting the pointers to special integer types – `intptr_t` and `uintptr_t`.

Linux manual page - `stdint.h(0p)`

The following type designates a signed integer type with the property that any valid pointer to void can be converted to this type, then converted back to a pointer to void, and the result will compare equal to the original pointer: `intptr_t`

The following type designates an unsigned integer type with the property that any valid pointer to void can be converted to this type, then converted back to a pointer to void, and the result will compare equal to the original pointer: `uintptr_t`




Figure 4.4: Visual example of an array-list based deque with a current front (fore) index allowing element traversal from first-to-last (fore-to-back) and last-to-first (back-to-fore).

4.1.3 Doubly-linked list of arrays

Combining the properties of both lists and arrays I managed create a deque with good locality and growth, look at Figure 4.4. Each linked list node is made of a previous next and elements_array member. Since it is this good is also gets described in more detail in the next section. Similar to the queue, this deque also relies on finding the perfect middle-ground between allocation speed (finding smaller memory chunks is faster) and locality (the bigger the memory array, the better the locality).

4.2 Doubly-linked list of arrays code

The structure consists of a child node structure representing a doubly-linked list node with a flexible array member for storing elements.

4.2.1 The infinite queue structure

The parent structure solely relies on the child head for element storage and traversal. The parent also stores the current front index to simplify front element access in `peek_front_ideque`. Next up is the `size` member storing individual element size, followed by `length`. `length` contains the number of elements in the deque. Lastly, the parent structure stores a pointer to a memory allocator structure for memory management.

`include/sequence/ideque.h`

```

/// @brief Doubly linked list node for deque structure.
struct infinite_deque_node {
    struct infinite_deque_node * next; // next sibling node
    struct infinite_deque_node * prev; // previous sibling node
    char elements[];                  // elements array with
    IDEQUE_CHUNK size
};

/// @brief Infinite deque data structure.
typedef struct infinite_deque {
    struct infinite_deque_node * head;
    size_t current, size, length; // current index, element size and
    structure length
    memory_s const * allocator;
} ideque_s;

```

4.2.2 Create deque

Creating a deque means only setting the size member to the size of the element type, and setting the allocator to the standard allocator pointer.

include/sequence/ideque.h

```

/// @brief Creates an empty structure.
/// @param size Size of a single element
/// @return Deque structure.
ideque_s create_ideque(size_t const size);

```

source/sequence/ideque.c

```

{
    \\ ...
    return (ideque_s) { .size = size, .allocator = allocator, };
}

```

4.2.3 Make deque

Making the structure is the same as creating it, but with the added option of using a custom memory allocator.

include/sequence/ideque.h

```

/// @brief Creates a custom empty structure.

```

```

/// @param size Size of a single element.
/// @param allocator Custom allocator structure.
/// @return Deque structure.
ideque_s make_ideque(size_t const size, memory_s const * const
    allocator);

```

The memory allocator simply gets set to the allocator member of the structure, while the size member is set the same way as in `create_ideque`.

source/sequence/ideque.c

```

{
    \\ ...
    return (ideque_s) { .size = size, .allocator = allocator, };
}

```

4.2.4 Destroy deque

Destroying the deque means freeing each element using the destroy function pointer and its arguments. The destroy function pointer can either do nothing, set bytes in the element to zero or free the element if it contains allocated memory.

include/sequence/ideque.h

```

/// @brief Destroys a structure and its elements, but makes it
    unusable.
/// @param deque Structure to destroy.
/// @param destroy Function pointer to destroy a single element.
/// @param ad Arguments for destroy function pointer.
void destroy_ideque(ideque_s * const deque, set_fn const destroy, void
    * const ad);

```

First a current node pointer is initialized to the head node, then a for loop sets a start index for the current index position for the special case that a node has elements not at the start of the array. All the other nodes have therefore elements starting at index zero, hence the start variable is set to zero after the first iteration.

source/sequence/ideque.c

```

{
    \\ ...
    struct infinite_deque_node * current = deque->head; // save head
    index as current node
    // for every node in deque list until size is not zero
    for (size_t start = deque->current, remaining = deque->length;
        remaining; start = 0) {

```

Following that, the index variable `i` is set to the start index, and a nested for loop iterates over the elements in the current node, destroying each element using the destroy function pointer. The loop continues until either all remaining elements have been destroyed or the end of the current node's elements array is reached. After destroying the elements, the remaining variable is updated to reflect the number of elements left to destroy.

source/sequence/ideque.c

```
    size_t i = start; // set i outside of nested for loop to save
    its value
    for (; i < (remaining + start) && i < IDEQUE_CHUNK; ++i) { //
    destroy each element in node
        destroy(current->elements + (i * deque->size), ad);
    }
    remaining -= (i - start); // subtract absolute value of i and
    start from deque's size
```

Lastly, the current node is saved as a temporary variable `temp` for later freeing. Next the current node pointer is updated to point to the next node in the list. Finally, `temp` is freed using the defined memory allocator's free function.

source/sequence/ideque.c

```
    struct infinite_deque_node * temp = current; // save current
    node as temporary to free later
    current = current->next; // go to next node

    deque->allocator->free(temp, deque->allocator->arg); // free
    temporary current node
}
```

All that is left is to zero-out the deque structure, which is done by using `memset`.

source/sequence/ideque.c

```
    // set everything to zero
    memset(deque, 0, sizeof(ideque_s));
}
```

4.2.5 Clear deque

Clearing the deque is similar to destroying it, but with the difference being that the structure remains usable. Usability means the deque can still support insertion, deletion and search.

include/sequence/ideque.h

```

/// @brief Clears a structure and destroys its elements, but remains
        usable.
/// @param deque Structure to destroy.
/// @param destroy Function pointer to destroy a single element.
/// @param ad Arguments for destroy function pointer.
void clear_ideque(ideque_s * const deque, set_fn const destroy, void *
        const ad);

```

source/sequence/ideque.c

```

{
    struct infinite_deque_node * current = deque->head; // save head
    index as current node
    // for every node in deque list until size is not zero
    for (size_t start = deque->current, remaining = deque->length;
    remaining; start = 0) {
        size_t i = start; // set i outside of nested for loop to save
        its value
        for (; i < (remaining + start) && i < IDEQUE_CHUNK; ++i) { //
        destroy each element in node
            destroy(current->elements + (i * deque->size), ad);
        }
        remaining -= (i - start); // subtract absolute value of i and
        start from deque's size

        struct infinite_deque_node * temp = current; // save current
        node as temporary to free later
        current = current->next; // go to next node

        deque->allocator->free(temp, deque->allocator->arg); // free
        temporary current node
    }

    deque->current = deque->length = 0;
    deque->head = NULL;
}

```

4.2.6 Copy deque

Copying is done on a per-element basis using iteration through each array element and node. The user is able to copy elements in any way they see fit, which also includes

deep/shallow copying, and reentrant copies.

include/sequence/ideque.h

```

/// @brief Creates a copy of a structure and all its elements.
/// @param deque Structure to copy.
/// @param copy Function pointer to create a deep/shallow copy of a
///         single element.
/// @param ac Arguments for copy function pointer.
/// @return Deque structure.
ideque_s copy_ideque(ideque_s const * const deque, copy_fn const copy,
    void * const ac);

```

We begin by creating a return replica deque structure which gets initialized with the same members as the deque to copy (except the head node pointer of course).

source/sequence/ideque.c

```

{
    // ...
    ideque_s replica = {
        .current = deque->current, .size = deque->size, .length =
        deque->length, .allocator = deque->allocator,
    };
}

```

First two node pointers are initialized. The first one is a simple current node pointer of the deque head to iterate over. Second is a double pointer to the memory location of the replica for simpler node update.

Next we initialize the for loop with a start index for the special case that the first node can have elements not at the zeroth array position. Continuing on, a remaining length variable is initialized to keep track of how many elements remain to copy.

source/sequence/ideque.c

```

struct infinite_deque_node const * current_deque = deque->head; //
save head index as current node
struct infinite_deque_node ** current_replica = &(replica.head);
for (size_t start = deque->current, remaining = deque->length;
remaining; start = 0) { // while remaining remains

```

Inside the loop, a new node is allocated and temporary set as node with proper memory size consisting of this formula:

$$sizeof_{node} + (chunk_{deque} \times size_{element})$$

A copied element counter *i* is set to start for the inner loop that iterates through each node's array until remaining length or IDEQUE_CHUNK is reached. Each element is

then copied into allocated node's array to the same index. Lastly, the remaining length is decremented by copied size $i - start$ to fix the start shift.

source/sequence/ideque.c

```
// allocate node for replica
struct infinite_deque_node * node = deque->allocator->alloc(
sizeof(struct infinite_deque_node) + (IDEQUE_CHUNK * deque->size),
deque->allocator->arg);
// ...

// copy each element from old to replica
size_t i = start;
for (; i < IDEQUE_CHUNK && i < remaining + start; ++i) { //
operate on each element in node
    copy(node->elements + (i * deque->size), current_deque->
elements + (i * deque->size), ac);
}
    remaining -= (i - start); // decrement remaining size by
operated size of current node
```

Lets see. The worst way to explain this code is to set the reason straight as to why `current_replica` is a double pointer. The answer is so that we can directly change the pointer value `current_replica` node in the deque struct or member in parent node struct instead of changing the value in the `current_replica` variable. First we make node circular to itself from its previous node, together with `current_replica`. That all is needed so the code remains general for both the cases of `current_replica` being either the head or body nodes.

Node's next member is set to head's previous pointer, because that guarantees circularity.

Node's previous member is reset to head's previous node, if `current_replica` is the head, then circularity is retained, otherwise the node's previous pointer gets corrected.

Head's previous node is set current node as the current node represents the last node added, which is where head's previous must point to. Overall, it is pretty hard code to understand just to avoid a couple if conditions.

source/sequence/ideque.c

```
// current_replica and node's prev will point to node
(*current_replica) = node->prev = node;
// node's next points to head (if current_replica is head then
node's next points to node)
node->next = replica.head;
// node's prev points to head's prev (if current_replica is
head then head/prev is node, else prev is previous node)
```

```

    node->prev = replica.head->prev;
    // head's prev points to node (if current_copy is head then
    head/node prev is node, else head prev is last node)
    replica.head->prev = node;

```

On a simple node, the last thing we do inside the loop is set `current_replica` and `current_deque` to the next pointer to pointer, and just pointer address, respectively. In the end, the properly initialized `replica` is returned.

source/sequence/ideque.c

```

    current_replica = &((*current_replica)->next); // go to next's
    node pointer
    current_deque = current_deque->next; // go to next node
}

return replica;
}

```

4.2.7 Is empty deque

Returns true if deque is empty.

include/sequence/ideque.h

```

/// @brief Checks if structure is empty.
/// @param deque Structure to check.
/// @return 'true' if empty, 'false' if not.
bool is_empty_ideque(ideque_s const * const deque);

```

source/sequence/ideque.c

```

{
    // ...
    return (deque->length == 0);
}

```

4.2.8 Enqueue deque front

Enqueues element to the front of the deque. The front is defined as the head node's current index element. The inserted element is provided as a pointer element parameter.

include/sequence/ideque.h

```

/// @brief Enqueues a single element to the front of the structure.
/// @param deque Structure to enqueue into.
/// @param element Element buffer to enqueue.
void enqueue_front_ideque(ideque_s * const deque, void const * const
    element);

```

First enqueue front checks if current index is zero. That is because if it is, then a new node needs to be created, pointers need to be redirected and saved to the structure.

source/sequence/ideque.c

```

{
    // ...
    if (!(deque->current)) { // if deque's previous current '
        underflows' in node array due to inserting element to front

```

The current index is changed to IDEQUE_CHUNK to later correct it to the last element in node's array during current decrement. Next a new node is allocated using the size of a singular node and chunk size multiplied with the size of a single element.

source/sequence/ideque.c

```

    deque->current = IDEQUE_CHUNK; // make current into list array
    chunk size to prevent future underflow

    // allocate node for replica
    struct infinite_deque_node * node = deque->allocator->alloc(
        sizeof(struct infinite_deque_node) + (IDEQUE_CHUNK * deque->size),
        deque->allocator->arg);
    // ...

```

The last thing done inside the if body is to check if a head exists (head != NULL). If head ain't NULL, node is properly redirected so its next pointer becomes head and its prev pointer becomes head's previous. Lastly, head's previous and head previous' next are set to node together with deque's head.

The else condition body makes the node circular and also sets deque's head to node. Which is way simpler to explain and understand than the if body statements.

source/sequence/ideque.c

```

    if (deque->head) { // if head exists
        node->next = deque->head; // node's next is head
        node->prev = deque->head->prev; // node's previous is tail
    }
    /head's previous

```

```

        deque->head = deque->head->prev = deque->head->prev->next
= node;
    } else { // else head does not exist
        deque->head = node->next = node->prev = node; // node's
next and previous is node and node becomes head
    }
}

```

The last part of the enqueue front function increments the deque's length and decrements current (to make it point to the last element in the array if it was zero (it circles back (to the beginning))). The element is memcopied to deque head element's current position.

source/sequence/ideque.c

```

    deque->length++; // increment size for new element insertion
    deque->current--; // if current was 0 then current will be '
IDEQUE_CHUNK - 1'
    memcpy(deque->head->elements + (deque->current * deque->size),
element, deque->size); // copy element into deque's head
}

```

4.2.9 Enqueue deque back

Enqueues element to the back of the deque. The back position is defined as an element located at the position $current + length - 1$. The element also gets saved as a element to that position.

include/sequence/ideque.h

```

/// @brief Enqueues a single element to the back of the structure.
/// @param deque Structure to enqueue into.
/// @param element Element buffer to enqueue.
void enqueue_back_ideque(ideque_s * const deque, void const * const
    element);

```

The next_index represents the next index where the element element gets saved relative to the node array. Relativity is achieved with IDEQUE_CHUNK modulo of $current + length$.

source/sequence/ideque.c

```

{
    // ...

```

```

size_t const next_index = ((deque->current + deque->length) %
IDEQUE_CHUNK);
if (!next_index) { // if next index to insert into is zero

```

Under the if condition a new node gets allocated.

source/sequence/ideque.c

```

    struct infinite_deque_node * node = deque->allocator->alloc(
sizeof(struct infinite_deque_node) + (IDEQUE_CHUNK * deque->size),
deque->allocator->arg);
    // ...

```

The nested if checks if head is NULL. If true, then the node is inserted just like in `enqueue_front_ideque`, but the node doesn't become the head node (it remains tailily, like... tail-esque). The else body is the same as the `enqueue_front`'s, since the head and tail are one and the same.

source/sequence/ideque.c

```

    if (deque->head) { // if head exists
        node->next = deque->head; // node's next is head
        node->prev = deque->head->prev; // node's previous is tail
/head's previous

        deque->head->prev = deque->head->prev->next = node; //
node is tail's next and head's previous, and node becomes tail
    } else { // else head does not exist
        deque->head = node->next = node->prev = node; // node's
next and previous is node and node becomes head
    }
}

```

Finally, all that is left is to only increment deque's length and memory copy to head previous', or tail's for short, element array at relative `next_index` index.

source/sequence/ideque.c

```

    deque->length++; // increment size for new element insertion
    memcpy(deque->head->prev->elements + (next_index * deque->size),
element, deque->size); // copy element into deque's tail
}

```

4.2.10 Dequeue deque front

The dequeue, unlike the enqueue, removes the front element from the deque, but like the enqueue, the dequeue works with the current index position. The buffer stores the removed element.

include/sequence/ideque.h

```
/// @brief Dequeues a single element from the front of the structure.
/// @param deque Structure to dequeue from.
/// @param buffer Element buffer to save dequeue.
void dequeue_front_ideque(ideque_s * const deque, void * const buffer);
```

Now the front-most element gets saved to the buffer and current increments, while the length decrements. Which reflects common deque behavior.

source/sequence/ideque.c

```
{
    // ...
    memcpy(buffer, deque->head->elements + (deque->current * deque->
size), deque->size);
    deque->current++; // increment current index since it's removing
from front (like a queue)
    deque->length--; // decrement deque's size since front elements
gets removed (like a queue)
```

Next up, if a node becomes empty and needs to be removed, an if condition checks if current was incremented to IDEQUE_CHUNK, which is an issue. Another way that a node may get removed is if length reaches zero. Back to the issue, if chunk size is reached than current index must be circled back.

source/sequence/ideque.c

```
// if deque's current index is equal to list chunk or deque is
empty free head node
if ((IDEQUE_CHUNK == deque->current) || !(deque->length)) {
```

This segment cuts off the head node through pointer list magic. it first saves the head so it doesn't roll off after the cut. Then head next's previous pointer points to tail and tail's next pointer points to head's next.

source/sequence/ideque.c

```
    struct infinite_deque_node * head = deque->head; // temporary
save pointer to head node
```

```

    head->next->prev = head->prev; // set head next's previous
    pointer to tail/head's previous
    head->prev->next = head->next; // set tail's next node to head
    's next node

```

Deque's head gets either set to its next or NULL if length reaches zero. *current* is set to zero for recircled, and lastly head node is memory freed.

source/sequence/ideque.c

```

    // if deque's size is zero then set head to NULL, else it's
    head's next node
    deque->head = deque->length ? deque->head->next : NULL;
    deque->current = 0; // reset current index to zero/beginning

    deque->allocator->free(head, deque->allocator->arg); // free
    temporary head node
}
}

```

4.2.11 Dequeue deque back

Dequeue back dequeues an element from the back of the deque and saves removed element into buffer. The back is defined as a position formula $current + (length - 1)$ modulo to single chunk size.

include/sequence/ideque.h

```

/// @brief Dequeues a single element from the back of the structure.
/// @param deque Structure to dequeue from.
/// @param buffer Element buffer to save dequeue.
void dequeue_back_ideque(ideque_s * const deque, void * const buffer);

```

Initially, the deque length is decreased to represent both removal and for the back position calculation. The position itself is represented by *back_index*, which is used to copy the back element into buffer.

source/sequence/ideque.c

```

{
    // ...
    deque->length--; // decrement only size since it's removing from
    the back
    size_t const back_index = (deque->current + deque->length) %
    IDEQUE_CHUNK; // calculate back element's index

```

```
memcpy(buffer, deque->head->prev->elements + (back_index * deque->
size), deque->size);
```

If deque is empty, that means only the head/tail node remain as one. Thus the only thing required is the free-up of the node's memory followed by setting the current and head deque struct members to zero/NULL for simpler future enqueues. If length is zero, then back_index may not be zero and it is a special 'remove node' case.

source/sequence/ideque.c

```
if (!deque->length) {
    deque->allocator->free(deque->head, deque->allocator->arg); //
    free head node

    deque->current = 0; // reset current index to 0 if deque is
    empty
    deque->head = NULL;
```

The first if is continued by an else if statement, which checks if back_index is zero, representing the removal of the last element in the node array. This requires freeing up the empty node and return unused memory back to the system. Since this node is the last node in the deque it is saved as a temporary tail for later freeing. Next the tail node is cut off with tail's previous node, and tail's previous node is made to point to deque's head instead of tail. The tail thus becomes lost in the deque, hence it was temporary saved, and all that is left is to free is with the allocator.

source/sequence/ideque.c

```
    } else if (!back_index) { // if current index at back is zero or
    deque is empty remove empty tail/head
        struct infinite_deque_node * tail = deque->head->prev; //
        temporary save pointer to head node

        deque->head->prev = tail->prev; // head's tail pointer equals
        tail's previous
        tail->prev->next = deque->head; // tail previous' next equals
        head

        deque->allocator->free(tail, deque->allocator->arg); // free
        temporary tail node
    }
}
```

4.2.12 Peek deque front

Peeks the front element in the deque at deque's current position.

include/sequence/ideque.h

```
/// @brief Peeks a single element from the front of the structure.
/// @param deque Structure to peek.
/// @param buffer Element buffer to save peek.
void peek_front_ideque(ideque_s const * const deque, void * const
    buffer);
```

The peeked element is saved into buffer without any removal.

source/sequence/ideque.c

```
{
    // ...
    memcpy(buffer, deque->head->elements + (deque->current * deque->
        size), deque->size);
}
```

4.2.13 Peek deque back

Peeks the back element in the deque at deque's $(current + (length - 1)) \bmod chunk$ position.

include/sequence/ideque.h

```
/// @brief Peeks a single element from the back of the structure.
/// @param deque Structure to peek.
/// @param buffer Element buffer to save peek.
void peek_back_ideque(ideque_s const * const deque, void * const
    buffer);
```

A `back_index` is calculated using the aforementioned formula and element is copied from tail (head's previous) position into buffer. The element is not removed.

source/sequence/ideque.c

```
{
    // ...
    // calculate back element's location/index and copy it into buffer
    size_t const back_index = (deque->current + deque->length - 1) %
        IDEQUE_CHUNK;
    memcpy(buffer, deque->head->prev->elements + (back_index * deque->
        size), deque->size);
}
```

```
}
```

4.2.14 Each deque front

Each element is iterated in front-to-back order using user-defined manage function pointer and its am arguments.

```
include/sequence/ideque.h
```

```
/// @brief Iterates over each element in structure starting from the
/// front.
/// @param deque Structure to iterate over.
/// @param manage Function pointer to operate on each element
/// reference using element size and arguments.
/// @param am Generic arguments to use in function pointer.
void each_front_ideque(ideque_s const * const deque, manage_fn const
    manage, void * const am);
```

A current node variable is set to deque's head in order to be later iterated inside a for loop. Next, a similar approach to destroy/clear deque is used to iterate through each node and its array elements until remaining unvisited elements count reaches zero. Also, a start variable is set to the current deque position for the special head node case where the current position may not lay at zero, while other nodes have the first element at zero.

```
source/sequence/ideque.c
```

```
{
    // ...
    struct infinite_deque_node * current = deque->head; // save head
    node as current pointer
    // while remaining size is not zero
    for (size_t start = deque->current, remaining = deque->length;
        remaining; start = 0) {
```

A nested for loop is iterating through each node array element starting at start. Inside the loop the manage function pointer is called with the current node element at i index and am arguments. The reason i is saved outside the for loop is to later subtract it from the remaining unvisited element count. Lastly, the current node is updated with its next sibling to continue the next iteration.

```
source/sequence/ideque.c
```

```
    size_t i = start;
    for (; i < (remaining + start) && i < IDEQUE_CHUNK; ++i) { //
        operate on each element in node
```

```

        // if operate returns false then terminate main function
        if (!manage(current->elements + (i * deque->size), am)) {
            return;
        }
    }

    remaining -= (i - start); // decrement remaining size by
    operated size of current node
    current = current->next; // go to next node
}
}

```

4.2.15 Each deque back

Each element is iterated in back-to-front order using user-defined manage function pointer and its am arguments.

include/sequence/ideque.h

```

/// @brief Iterates over each element in structure starting from the
/// back.
/// @param deque Structure to iterate over.
/// @param manage Function pointer to operate on each element
/// reference using element size and arguments.
/// @param am Generic arguments to use in function pointer.
void each_back_ideque(ideque_s const * const deque, manage_fn const
    manage, void * const am);

```

A current node is defined as the deque head, then a remaining size is defined as the deque's length to represent the count of unvisited elements. Next a last index variable is defined to access the last element in the whole structure (not relative to chunk).

source/sequence/ideque.c

```

{
    // ...
    struct infinite_deque_node * current = deque->head; // save head
    node as current pointer
    for (size_t remaining = deque->length, last = deque->current +
        deque->length - 1; remaining;) {

```

Inside the for loop, the current node is updated to its previous node to both start at the tail node and to iterate through the deque in reverse. Later a next_index is defined to access the node's elements relative to IDEQUE_CHUNK size. An i index is defined outside

another for loop which iterates over each element in node array in reverse. The inner for loop checks if *i* is less than or equal to *node_index* (yes, equal to, because *node_index* is a zero-based index *i* must be equal to, and not a length). The for loop also checks if any remaining elements exist, at the end, the loop increments *i* and decrements *remaining*.

source/sequence/ideque.c

```

    current = current->prev; // fist go to tail and other previous
    nodes

    size_t const node_index = last % IDEQUE_CHUNK; // calculate
    last element index in node
    size_t i = 0; // save number of operated elements in node

```

Inside the inner for loop, the manage function pointer is called for each element in reverse of *i*, as *node_index - i*.

source/sequence/ideque.c

```

    for (; i ≤ node_index && remaining; ++i, remaining--) { //
    until node index is reached and elements remain
        // manage on reversed i to start from last element in node
        if (!manage(current->elements + ((node_index - i) * deque
        ->size), am)) {
            return;
        }
    }
    last -= i; // may underflow
}

```

4.2.16 Apply deque

Turns the deque into an array which the user can manipulate, like sort the elements for example, and the 'changed' array is saved back into the deque. Can also be treated like a continuous buffer for printing or more cache friendly elements access as oppose to the each functions.

include/sequence/ideque.h

```

/// @brief Apply each element in structure into an array to manage.
/// @param deque Structure to map.
/// @param process Function pointer to manage array of elements using
    structure length, element size and arguments.
/// @param ap Generic arguments to use in function pointer.

```

```
void apply_ideque(ideque_s const * const deque, process_fn const
    process, void * const ap);
```

Since the deque does not save all element into as a continuous array, the function first allocates an `elements_array` of $(length \times size)$ size.

source/sequence/ideque.c

```
{
    // ...
    // create elements array to temporary save doubly linked list
    elements into straight array
    char * elements_array = deque->allocator->alloc(deque->length *
        deque->size, deque->allocator->arg);
```

Next the temporary `elements_array` is initialized with the elements from the deque using the same front-to-back iteration for loop approach, explained in `destroy_ideque`, thus also in `clear_ideque`, and `each_front_ideque`. An index variable is used to properly save elements into the `elements_array` in zero-based index order.

source/sequence/ideque.c

```
    // index increments after getting each element from node
    // remaining count of elements is calculated
    // copying starts from head
    size_t index = 0, remaining = deque->length;
    struct infinite_deque_node * current = deque->head;
    // set start to current and then begin from 0 for all the other
    elements in array node
    for (size_t start = deque->current; remaining; start = 0) {
        size_t i = start; // temporary save iterator outside for later
        subtraction

        // copy elements starting from start until either 'relative
        remaining' or 'array node' size is reached
        for (; i < (remaining + start) && i < IDEQUE_CHUNK; ++i) {
            memcpy(elements_array + (index * deque->size), current->
elements + (i * deque->size), deque->size);
            index++;
        }

        remaining -= (i - start); // subtract copied count from
remaining size
        current = current->next; // go to next node
    }
```

After the `elements_array` is initialized, the process function pointer is called and the array gets processed by the user function together with `length` and `ap` arguments as parameters.

source/sequence/ideque.c

```
process(elements_array, deque->length, ap);
```

Then the index parameter is reset to zero and the `elements_array` elements are saved back into the deque structure using the same deque iteration algorithm.

source/sequence/ideque.c

```
// reset index and remaining and copy back elements from array
// into each node array
index = 0, remaining = deque->length, current = deque->head;
// set start to current and then begin from 0 for all the other
// elements in array node
for (size_t start = deque->current; remaining; start = 0) {
    size_t i = start; // temporary save iterator outside for later
    subtraction

    // copy elements starting from start until either 'relative
    // remaining' or 'array node' size is reached
    for (; i < (remaining + start) && i < IDEQUE_CHUNK; ++i) {
        memcpy(current->elements + (i * deque->size),
        elements_array + (index * deque->size), deque->size);
        index++;
    }

    remaining -= (i - start); // subtract copied count from
    remaining size
    current = current->next; // go to next node
}
```

Lastly, the `elements_array` gets deallocated.

source/sequence/ideque.c

```
deque->allocator->free(elements_array, deque->allocator->arg);
}
```

4.3 Example usages

Here are some example usages of the infinite deque. Each example show a deque of integers performing some of the functionalities mentioned in the previous section.

4.3.1 Creating, cleaning and destroying a deque

A simple program example which defines a destroy function pointer, creates an integer deque, clears and destroys the structure with the function pointer.

```
#include <sequence/ideque.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

int main(void) {
    ideque_s deque = create_ideque(sizeof(int));

    // some operations

    clear_ideque(&deque, intdst, NULL);

    // some new operations

    destroy_ideque(&deque, intdst, NULL);

    return 0;
}
```

4.3.2 Enqueueing, peeking and dequeuing elements front and back

Example showing the usage of front and back enqueueing, dequeuing and peeking elements.

```
#include <stdio.h>
#include <sequence/ideque.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}
```

```

int main(void) {
    ideque_s deque = create_ideque(sizeof(int));

    // enqueue front and back 10 elements
    // 9, 8, 7, 6, 5, 4, 3, 2, 1, 0, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
    for(int i = 0; i < 10; i++) {
        enqueue_front_ideque(&deque, &i);
        enqueue_back_ideque(&deque, &i);
    }

    // see if front and back elements are 9
    int a = -1, b = -1;
    peek_front_ideque(&deque, &a);
    peek_back_ideque(&deque, &b);
    printf("%d = %d\n\n", a, b);

    // dequeue and print all elements
    while(!is_empty_ideque(&deque)) {
        dequeue_front_ideque(&deque, &a);
        dequeue_back_ideque(&deque, &b);
        printf("%d = %d\n", a, b);
    }

    destroy_ideque(&deque, intdst, NULL);

    return 0;
}

```

4.3.3 Iterating and sorting elements

Code shows how iteration works with printing each element in the deque in front and back order. After that, sorting is shown using `qsort`.

```

#include <stdio.h>
#include <stdlib.h>
#include <sequence/ideque.h>

void intdst(void * const element, void * null) {
    (void)(null);
    (*(int*)element) = 0; // to set to 0, or just use (void)(element);
}

```

```

int intrcmp(void * const a, void * const b) {
    return (*(int*)b) - (*(int*)a); // reverse comparison
}

bool intprt(void * const element, void * format) {
    printf(format, (*(int*)element));
    return true;
}

void intsrt(void * const array, size_t const length, void * const
    compare) {
    qsort(array, length, sizeof(int), compare);
}

int main(void) {
    ideque_s deque = create_ideque(sizeof(int));

    // enqueue 10 elements
    // 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
    for(int i = 0; i < 10; i++) {
        enqueue_back_ideque(&deque, &i);
    }

    // print each element from start to end in order
    // prints: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
    each_front_ideque(&deque, intprt, "%d ");
    putc('\n', stdout);
    // prints: 9, 8, 7, 6, 5, 4, 3, 2, 1, 0
    each_back_ideque(&deque, intprt, "%d ");

    // sort elements in reverse
    // YOU SHOULD NOT CAST A FUNCTION POINTERS INTO VOID*
    // ARGUMENTS, THIS IS JUST A SILLY EXAMPLE, INSTEAD PUT
    // THE FUNCTION POINTER INSIDE A STRUCT AND PUT THE
    // INITIALIZED STRUCT POINTER ITSELF AS AN ARGUMENT
    apply_ideque(&deque, intsrt, intrcmp);
    putc('\n', stdout);

    // print each element from start to end in reverse
    // prints: 9, 8, 7, 6, 5, 4, 3, 2, 1, 0
    each_front_ideque(&deque, intprt, "%d ");

    destroy_ideque(&deque, intdst, NULL);
}

```

```
    return 0;  
}
```

4.4 On a final note...

C++'s main implementation of the deque [1], from which stack and queue inherit their functionalities. By the way, since I haven't quoted wikipedia yet, here is an insightful wiki article about deques [2].

Further reading

- [1] cppreference.com. std::deque. <https://en.cppreference.com/cpp/container/deque>. (Visited on 07/21/2026).
- [2] Wikipedia. Double-ended queue. https://en.wikipedia.org/wiki/Double-ended_queue. (Visited on 07/21/2026).

Part II

Linked lists

It is probably wrong to call this book's linked list part implementations 'linked lists', at least in the traditional C sense. These structures are more like 'bad' array, but not like array lists, where element retrieval is constant. No, no, no, they are linked, not by a memory address, but by array indexes.

Chapter 5

Straight list

We have a video, since videos can be simplified as frames, or images, in sequence. Our access to all elements mustn't be restricted, so we can't use Restricted sequential data structures.

One solution would be to use arrays, but that is where a problem might emerge. The array would save each frame inside it in continuous memory, so if a one hour video consists of $30_{fps} \times 60_{sec} \times 60_{min} = 108,000$ frames in total, and each frame is an image with $4_{RGBA} \times 1920 \times 1080 = 8,294,400$ bytes (8.3MB), the total is $108,000_{frames} \times 8,294,400_{bytes} = 895,795,200,000$ or ~896GB, wait, what, i-its actually this much? Thankfully we have video-compression. Now, it is pointless to use an array, but if we have a linked list with one 8.3MB (plus 8 bytes for a pointer) frame each scattered throughout our memory, we could maybe at least store it in memory holes.

Another solution is to treat each frame as a pointer and store them in an array of $108,000_{frames} \times 8B_{pointer} = 864,000$ byte size. The frames are scattered throughout the memory if we allocate space for each separately.

Using linked lists as a linear representation for each frame in a video gives us the ability to remove (cut out) parts of the video. Another feature is inserting a set of frames at any point in the video (combine).

A straight list is a special name I came up with to describe a singly-linked acyclic list. Adding and removing frames is easiest when done from the beginning, or head. That is because the to get to any other part of the list, a saved link variable must traverse through the entire structure to reach the end. This is most similar to how the Stack adds and retrieves the last inserted element first in the fastest way possible.

5.1 Implementations of a straight list

The straight list can have either one or many forms of realizations in C, depending on how far-stretched the definition of a singly-linked acyclic list is. But all can agree that

		PACKING	
		sparse	dense
ORDERING	pseudorandom	linked list straight list + circular list sparse array*	swapback array double list
	sequential	sparse array*	array list

Figure 5.1: Classification of list-esque data structures. The ones in bold contain this part's list implementations. PACKING refers to how the data is stored inside the array - sparse means holes, dense no holes. ORDERING is based on array, not list, order relative to all memory or chunks of it.

*sparse arrays are usually implemented as pseudorandomly ordered arrays, but can also be implemented as sequentially ordered arrays with boolean indicator arrays.

the base implementation consists of a structure called a node, which contains a value and a reference. The reference, in C a pointer, references either the next node in the list or a special reference symbolizing the end of the list (NULL).

Linux manual page - NULL(3const)

NULL represents a null pointer constant, that is, a pointer that does not point to anything.

Other forms of lists, like in this chapter, represent lists with arrays and are more concerned with how nodes are stored inside arrays. The issues they solve center around value packing inside arrays (sparse and dense arrays), and array node order based on insertions (sequential and pseudorandom ordering). The classic linked list can also be viewed from an array perspective, since the used memory can be represented as a one-dimensional array, see Figure 5.1.

This chapter will only cover the classic linked list, sparse array and straight list.




Figure 5.2: Linked list example with six nodes laid out in memory. The order in memory is not sequential and the nodes themselves are also sparsely packed.

5.1.1 A bread of crumbs - the linked list

The ordering of elements in a linked list (in memory) depends on the function which retrieves an uninitialized point in memory. Mostly that is done using the `malloc()`, short for [m]emory-[alloc]ation, function. Most lectures teach linked lists in C with the help of `malloc()`. With `malloc()`, memory is given to you in a way that is not guaranteed to be continuous, so the order of elements in memory is not guaranteed to be the same as the order of elements in the list, see Figure 5.2. Ideally, the order of elements in memory should be the same as in the list, mainly to avoid jumping back and forth in memory, which was visited either directly or in the area. On the other hand the memory can be too far away, so visiting it and returning will simply add, spoiler alert - pointless overhead.

The simplest solution would be to not use `malloc`. As long as the list has no references, only data, and is only being inserted into, we can use an array using `malloc()`, if the array is short, we can use `realloc()` local memory, or even global memory!!! So what if we are just stuck with a finite array to store data, the `malloc()` school class enemy is no more. So what to do if we wanna allow element removal?

5.1.2 A slice-off emmental - the sparse array

The sparse array allows element removal, but under some conditions:

1. The sparse array should insert the next element to an empty location as fast as possible.

	0	1	2	3	4	5
						
	3	0	2	1	0	1
ID	P	R	N	B	Q	K

Figure 5.3: A special value sparse array example using the count of playable chess pieces as an indicator that the element, i.e. chess piece, is not available. We got lucky that the chess piece count of zero also represents the non-existence of a piece on the board. The bottom letters are short for (P)awn, (R)ook, K(N)ight (rhymes with (N)ight), (B)ishop, (Q)ueen, and (K)ing. The top part above the letters represents the chess pieces as objects with attributes - like position on the board or usage count, values represent the count and the images represent image data to be rendered on the screen (notice how all the data is tightly packed, but the holes are creating gaps). The topmost part are array indexes, kind of like ids in a database.

2. If an element is removed, the array must indicate that the place the element was saved is empty, or to be more precise, inaccessible.

Solution number one must therefore save a count of all insertion calls, not just the current number of elements in the sparse array. This is because the next element must be inserted at the rightmost position to guarantee constant time insertion. To add a new piece into Figure 5.3, we could replace an empty piece through iterating the array and finding a free space, but we would lose the information of the piece not existing anymore. So if a new piece appears in chess, then we can only append it to the right. Removal is simply the absence of a piece, so to remove one, index to the piece and set its count to zero - constant time. Iterating for non-empty pieces can be slower due to caching, since some pieces may be separated from others like islands in the ocean.

The one point fifth solution requires another array of Booleans or even bits to check if a space is empty - true if filled, false if empty. This solution is referred to as a bit array, see Figure 5.4. Insertion and deletion are similar to solution number one, but instead of storing a count of used pieces (with count having multiple bytes), we store just one lookup bit. The disadvantage is that bit manipulation and bit storage are a little bit more difficult to wrap your head around. Also iterating through the bit array is slightly faster since finding the next bit in a bit array is faster than finding the next non-empty piece in a special value sparse array as you load 8x more bits than bytes.

	0	1	2	3	4	5	6	7
								
	1	0	0	0	0	0	0	1
ID	P0	P1	P2	P3	P4	P5	P6	P7

	8	9	10	11	12	13	14	15
								
	0	1	1	0	1	0	1	0
ID	R0	N0	B0	Q	K	B1	N1	R1

Figure 5.4: A bit array example using bit to show which pieces are playable. The bottom black letters are short for (P)awn, (R)ook, K(N)ight (still rhymes with (N)ight), (B)ishop, (Q)ueen, and (K)ing, the numbers next to some letters are used to differentiate between the same piece type, like two rooks, so rook 0 and rook 1. The smaller values above the black letters represent the set bits, with zero meaning the chess piece is dead, and one meaning the chess piece is alive on the board. Above them are chess piece objects and their data. The topmost part are array indexes, kind of like ids in a database. Notice the huge gap between the first and last pawn, which may have an impact on performance if the data is not cached.




Figure 5.5: A sparse set example showing which pieces are playable. The bottom black letters are short for (P)awn, (R)ook, K(N)ight (still rhymes with (N)ight), (B)ishop, (Q)ueen, and (K)ing, the numbers next to some letters are used to differentiate between the same piece type, like two rooks, so rook 0 and rook 1. The smaller values below the black letters represent the sparse array indexes pointing to its dense array sibling. The '-' represents the empty spaces to lookup if a piece is not present anymore. The dense array is above with indexes pointing back to the sparse array and also contains a densely packed array of chess piece objects.

Solution number two has two common solutions, either a special value is used to indicate that the element is empty, or another array is used to tell where to skip to the next valid element. The first solution still forces you to check iteratively if the memory space is valid, which can be processor intensive if elements are separated by oceans of special empty values. And the worst part, we lose the special value, that means we won't be able to use it if needed. This solution is also called a Sparse array.

The second solution is more efficient, but requires more memory to store the second array and more logic. The empty spaces are also called holes, so if there is a hole, then why not just fill it up?

5.1.3 A stack of butter - the straight list

Holes are a problem in sparse arrays, so the solution is to fill them. To fill them, though, we must track them. To track them, we must store them, and the simplest way to store and retrieve them is in a stack. The stack is a data structure that allows you to store and retrieve elements in a last-in-first-out (LIFO) manner. One way to implement the stack is to use a separate array to store the hole indexes. The straight list way of implementing

a stack is to use the same list, more specifically, add one more head reference to the list, which will point to the first hole in the list. The straight list structure will serve as both the linked list for elements and the stack for empty spaces. With this implementation we lose the sequential array ordering of the sparse list, but the sparsity problem gets minimized (not solved). To insert an element, we first check if the stack is empty, if it is, we insert the element at the end of the list, otherwise we pop the top of the stack and insert the element at that hole index.

The structure works like a dynamic array, the straight list must expand when elements are inserted and capacity is hit, and shrink if a smaller capacity is available. The biggest liability is thus shrinking, because if the stack has holes, we can't shrink the list array via reallocation. We must allocate a new array, copy the elements over in list order, and free the old array. The only benefit of this is that the shrunken structure has the same list and array element order - the element packing is dense and ordering is sequential.

A way to remove excess holes after each removal is to only push indexes of holes that are not at the end of the underlying array. When the last element from the list is removed, we can also reset the stack back to empty, by just setting its head reference to NULL.